

Approximate nearest neighbours search

Stanislav Protasov

Agenda

- ANNS (not ANNs)
 - Clustering and IVF
 - Proximity graphs (NSW, HNSW)

Before we start...

What's wrong with inverted index in terms of data structure?

Do you know the difference:

- $O(f(N))$,
- $O_A(f(N))$,
- $E(f(N))$?

Approximate Nearest Neighbours Search

Approximation for k-NN search

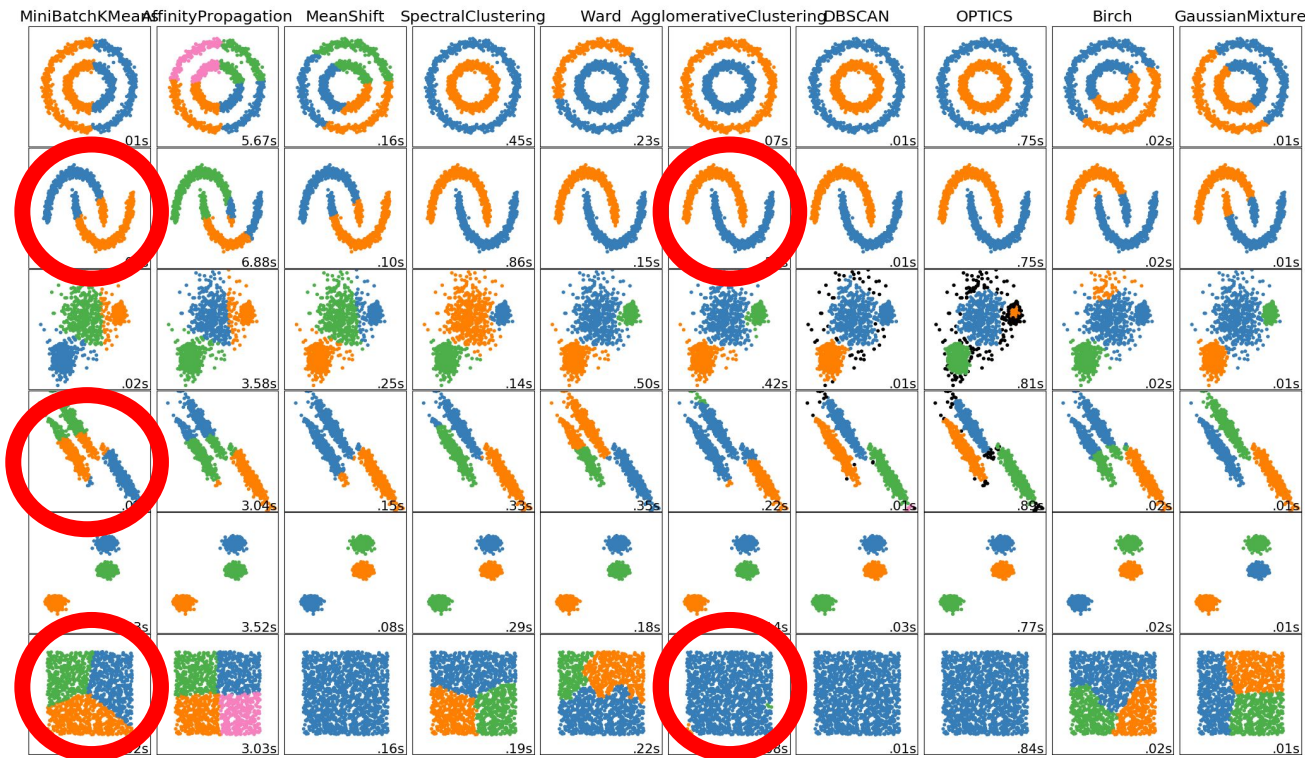
1. Pre-select $k \cdot c$ elements from approximate neighbourhood (~pre-ranking set). * Practical example: **Recall@1000 ≥ 0.875**
2. Then select and re-rank relevant ones.

How to:

- Locality sensitive hashing
- **Search trees and supporting data structures**
- **Vector compression, clustering, inverted indexing**
- **Proximity graphs**

Hierarchical clustering and Inverted index revised


How clustering approaches differ?



Linkage criteria

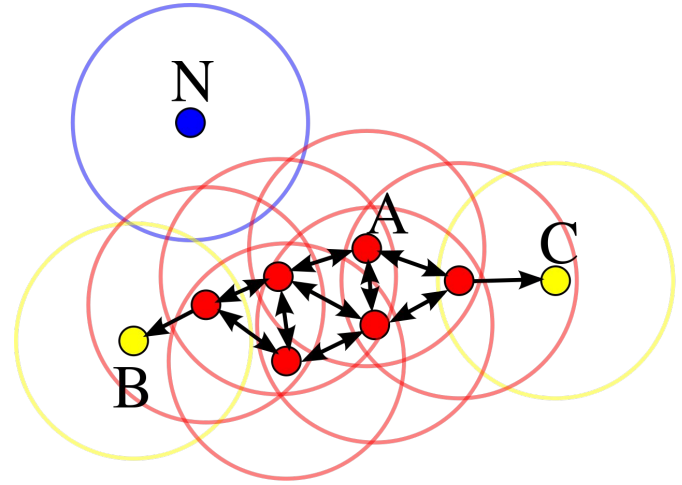
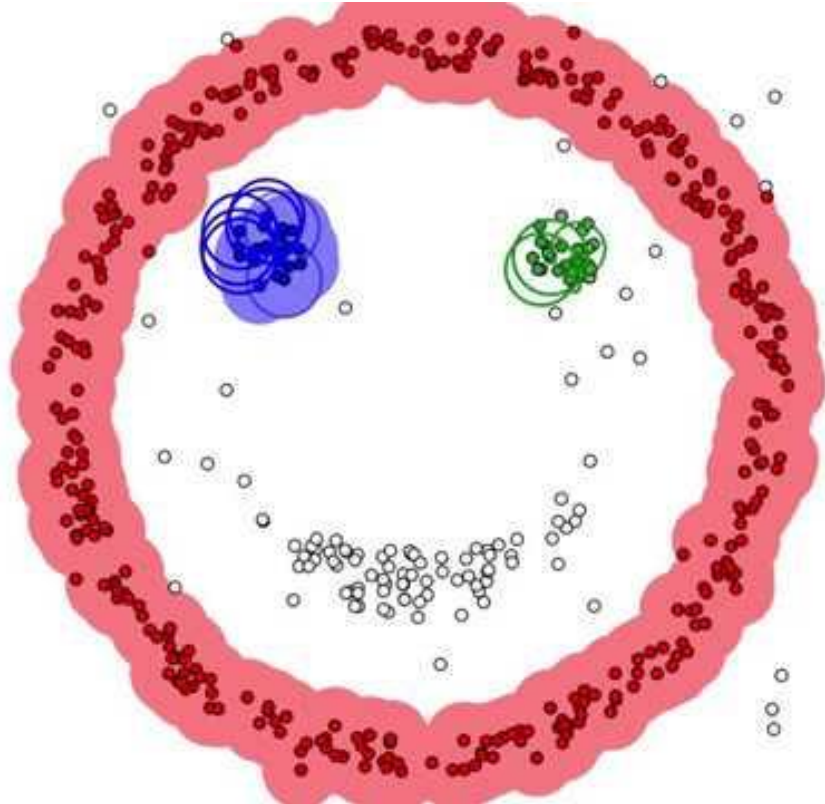
- **Single linkage** (smallest distance) ~ DBSCAN
- Complete linkage (maximum distance)
- Minimum energy (variance grows slowly in we merge)
- Average distance and **centroid-based approaches** — kMeans

K-Means



DBScan

Density-based spatial clustering of applications with noise.



Why do we cluster?

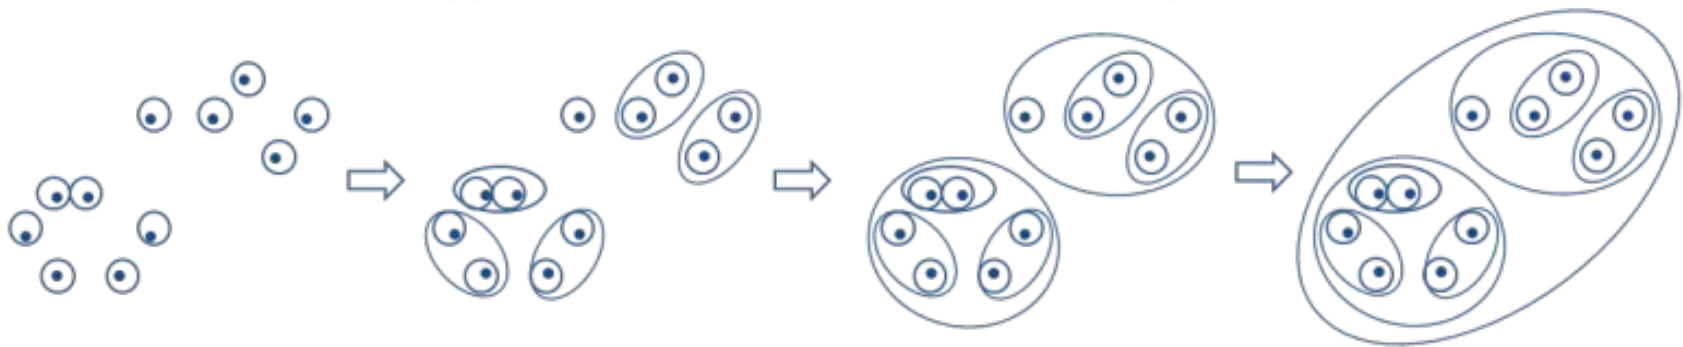
For a flat list we run $O(N)$ comparisons to find kNN

For \sqrt{N} similar^{*} clusters we can pick one closest^{**}
for $O(\sqrt{N})$ and find ***k*** NNs^{***} in $O(\sqrt{N})$.


For two layers of $\sqrt[3]{N}$...

How do we cluster?

Agglomerative Hierarchical Clustering

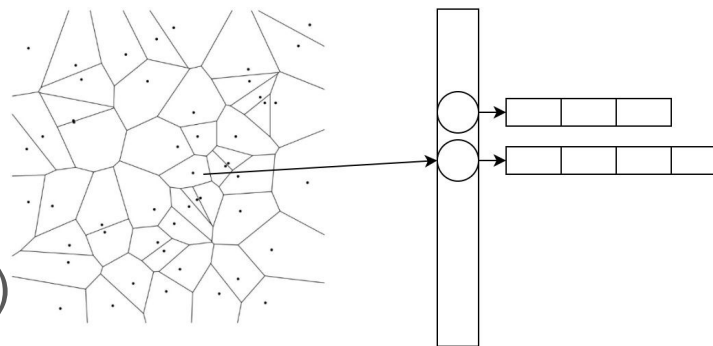
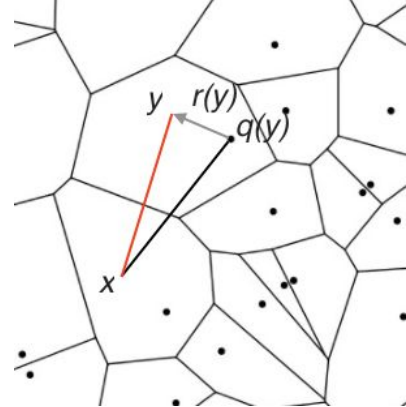


Divisive Hierarchical Clustering

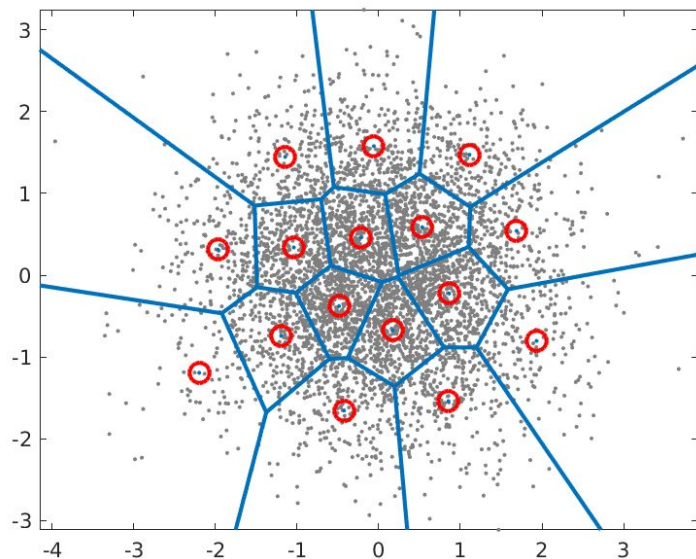


Revised IVF. [FAISS](#) (Facebook AI similarity search)


- Uses [Voronoi diagram](#) clusters (kMeans).
Vectors are approximated with **centroids (VQ)**
And compressed with
scalar quantization (**SQ**)
- Build **inverted index** for points in clusters
- Vector compression: product quantizer (**PQ**)
 - Split R^{128} into 8 groups of 16 floats
 - Perform 256-means clustering of these “sub-vectors” and encode with 1 byte each



Vector Quantization and Product Quantization



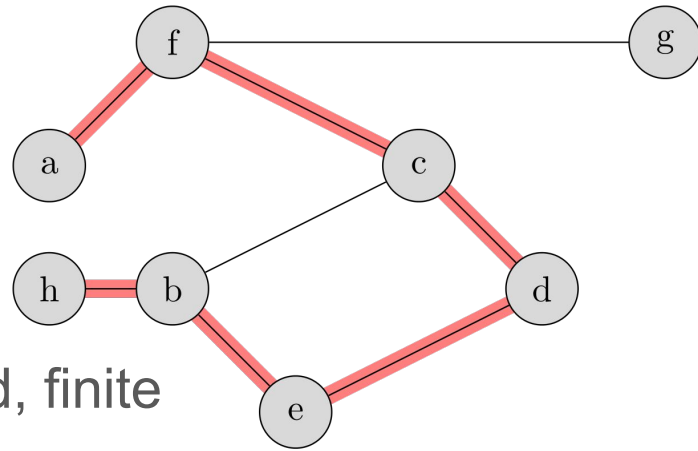
<https://wiki.aalto.fi/pages/viewpage.action?pageId=149883153>



<https://arbabenko.github.io/MultiIndex/>

Approach #2. Proximity graphs

Graphs cheat sheet



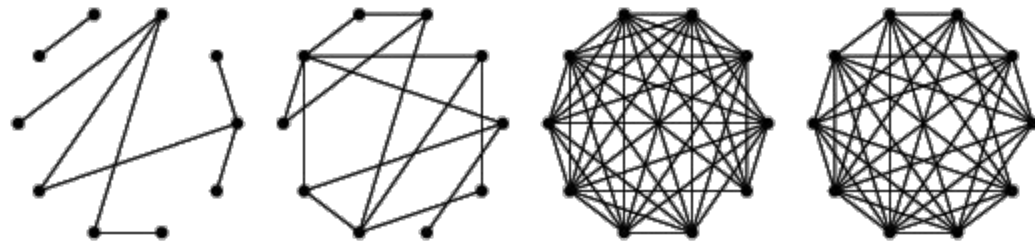
Graph - $G = (V, E)$, can be weighted, directed, finite

[Simple] **path** - sequence of vertices and edges

Degree of vertex - number of incident edges

Graph diameter - longest shortest path between a pair of vertices

Random graph



Some random process (uniform, Gaussian, ...) generates edges.

Almost every graph in the world. *Previously* considered as a model for social networks.

Small average shortest path - which is **good** for **search**.

Small clustering coefficient (defines how close are neighborhoods to cliques) - which is **bad** for **NN search**.

$$C(v) = \frac{e(v)}{\deg(v) (\deg(v) - 1) / 2}$$

$$\tilde{C} = \frac{1}{N} \sum_{i=1}^N C(i)$$

Regular graphs

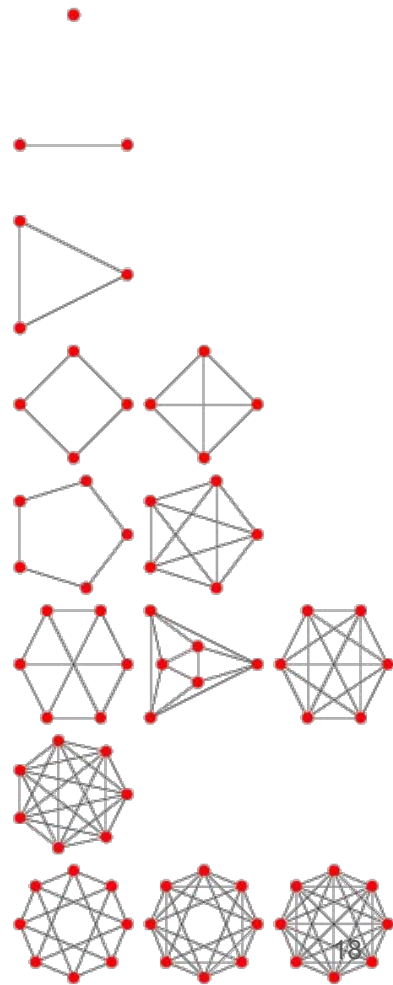
K-regular graph is a graph with $\deg(v) = K$ for any v .

Used to model big homogeneous networks.

Can also be random (as there are multiple K-regular graphs on the same size)

Big diameter - which is **bad** for **search**

Big clustering coefficient - which is **good** for **NN search**



Small World experiment by Stanley Milgram, 1967

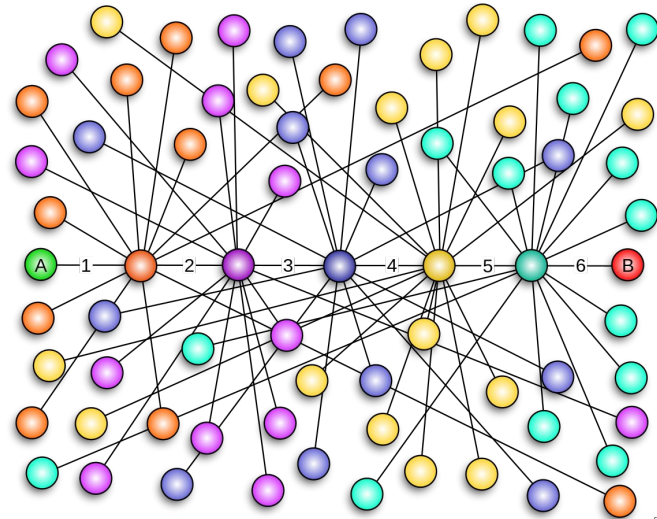
Initially it was considered, that social graph is kind of regular.

Experiment discovered (even with some questions to method) that even graph is **highly clustered, average path length is small.**

Was a basis for 6 handshakes rule.

New type of graphs was suggests:

small world networks.



Small world network

Most vertices are not neighbours (small degree means *sparse* graph).

Nevertheless, small number of hops needed to reach any other node.

Typical path length L between 2 random nodes (of N): $L \propto \log N$

Many real world networks are like this: internet, wiki, social graphs, power grids, brain cells. Although not all real networks like SW: many-generation networks, classmate graphs.

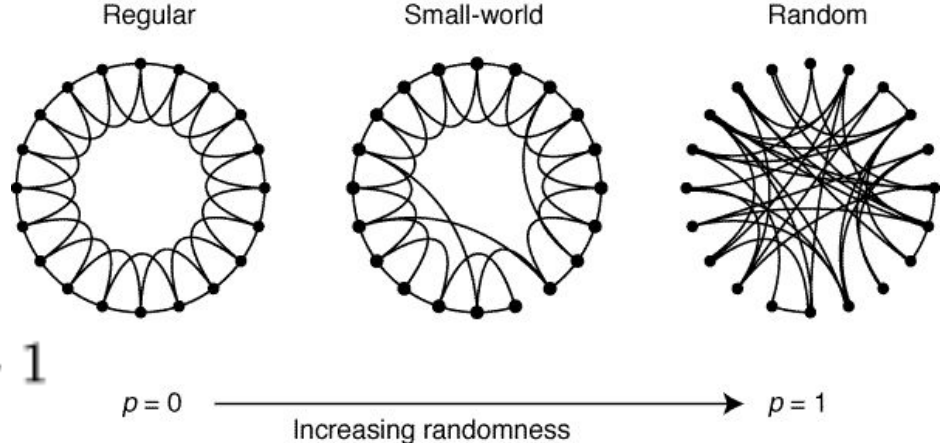
Watts–Strogatz model and Kleinberg model are how we describe and build SW networks

Watts–Strogatz model

Given N nodes and K -“regularity”
(average degree K) $N \gg K \gg \ln N \gg 1$

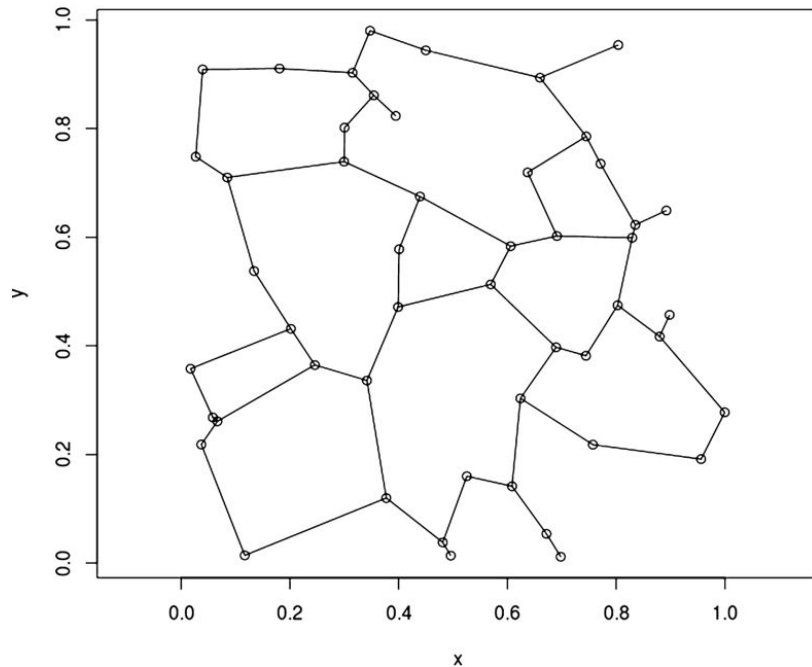
Given parameter p from $[0, 1]$.

- 1) Construct a regular ring lattice.
- 2) take every edge connecting **vertex** to its $K/2$ **rightmost neighbors**, and rewire it with probability p . Rewiring is done by **replacing destination** with vertex k (chosen **uniformly** at random from all possible nodes while avoiding self-loops and duplication).



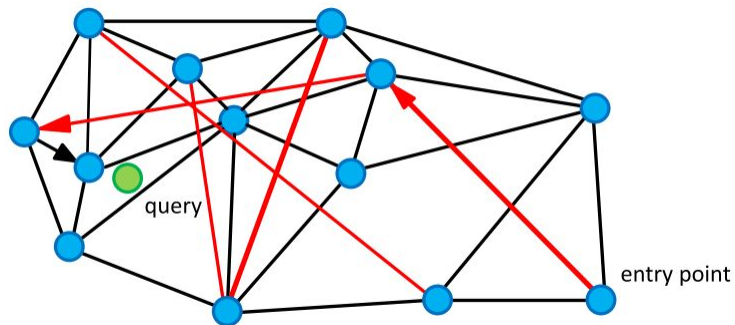
Proximity graphs

A proximity graph is simply a graph in which two vertices are connected by an edge if and only if the vertices satisfy particular geometric requirements.



Navigable small world networks

Idea is similar to [skip-lists](#).



We can also measure **distance** (e.g. dot product, Euclidian, L_k -norm, Humming, Levenstein, ...) between *query* and *current vertex*. Originally *Delaunay graph* needed to converge for exact search, but ANNS allows other small-world graphs.

Building:

1. One-by-one insertion via kNN search. Distant edges are created in the beginning.

Search:

1. Perform greedy search. Move to the neighbour vertex **closest to query**
2. Update NN set on each step until it converges

Hierarchical navigable small world ([github](#))


Layer 0 holds complete NSW network

Ideas:

- Better **start search** from a node with **high degree**
- Higher layer has longer links (skip-list!)
- Decrease layer size exponentially

Highlight:

- 1) **search procedure requires only $\text{dist}(u, v)$ function**
- 2) **No embedding, hyperplanes, centroids of whatever needed**



To read

[An Introduction to Proximity Graphs](#)

[Efficient and robust approximate nearest neighbor search using Hierarchical Navigable Small World graphs](#)

[How can proximity graphs benefit in classification and hybrid memory](#)