

# Lecture#1: Introduction to Distributed Systems

S. M. Ahsan Kazmi

# Outline

- Introduction
  - What, why, advantages & challenges?
  - Basics of distributed systems

# What are Distributed System?

- A distributed system:
  - Multiple connected CPUs working together
  - A collection of independent computers that appears to its users as a single coherent system
- Examples: parallel machines, networked machines

# Why Distributed Systems?

- Many systems that we use on a daily basis are distributed
  - World wide web, Google
  - Amazon.com
  - Peer-to-peer file sharing systems
  - Grid and cluster computing
  - Modern networked computers
- Useful to understand how such real-world systems work
- Course covers basic principles for designing distributed systems

# Advantages and Disadvantages

- Advantages
  - Communication and resource sharing possible
  - Economics – price-performance ratio
  - Reliability
  - Scalability
  - Potential for incremental growth
- Disadvantages
  - Distribution-aware Programming Languages, OSs and applications
  - Network connectivity essential
  - Security and privacy

# Goals of Distributed Systems

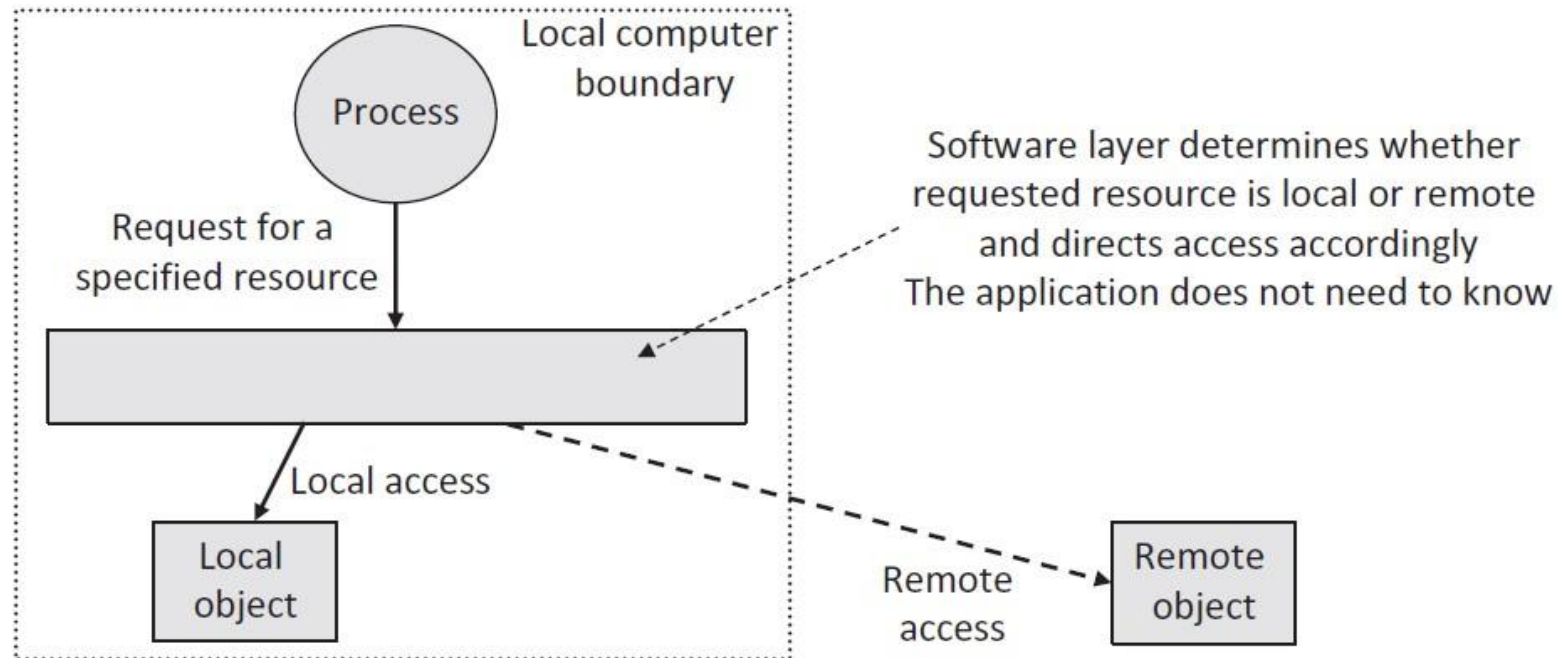
- Making resource accessible or sharable
- Distribution transparency
- Openness
- Scalability

# Transparency in a Distributed System

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource may be shared by several competitive users
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource
Different forms of transparency in a distributed system.	

# Access Transparency

- Access transparency requires that objects (this includes resources and services) are accessed with the **same operations** regardless of whether they are **local or remote**.





# Location Transparency

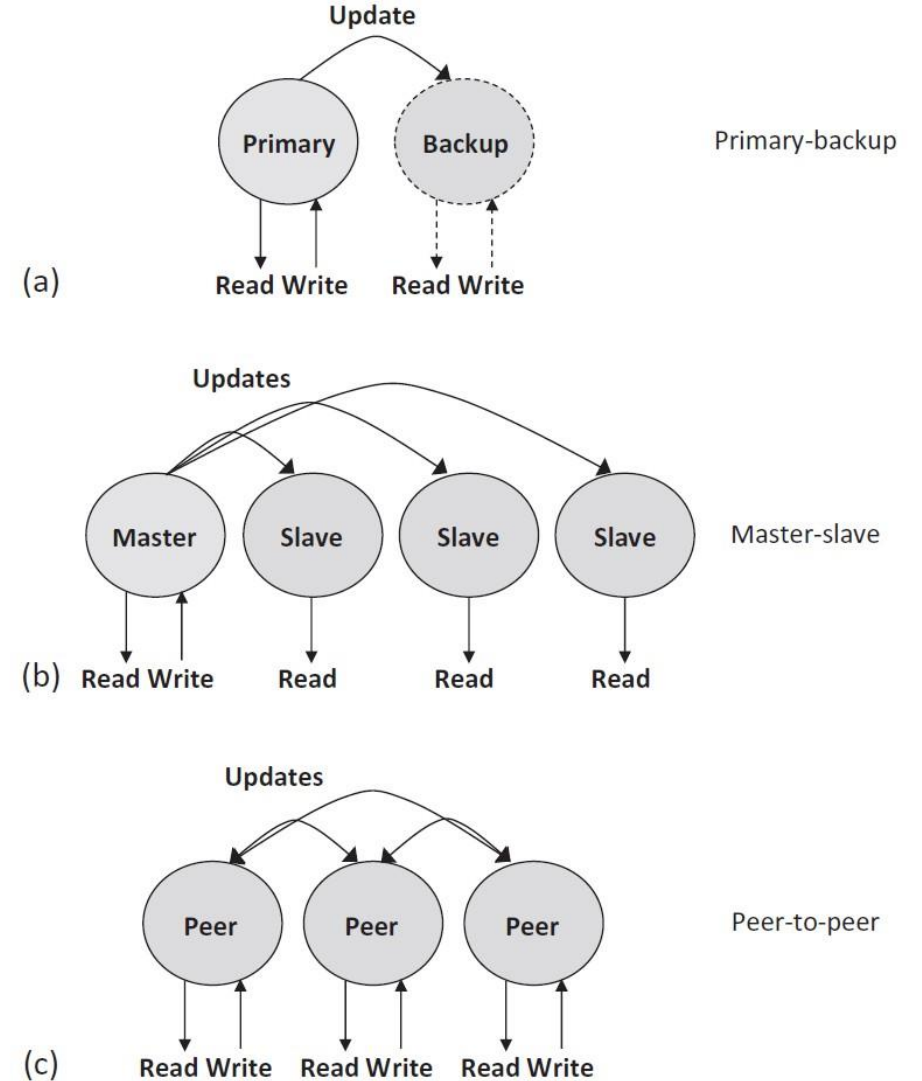
- Location transparency is the **ability to access** objects without the **knowledge of their location**
- The provision of location transparency is often achieved through the use of special services whose role is to perform a **mapping** between a **resource's name and its address**
- Resource virtualization (“access transparency”) – also provides location transparency

# Relocation Transparency

- Hide that an object **may be moved** to another location **while in use**
- The system may decide to move data to another location for performance optimization reasons
- Example: System automatically moves an object to another location due to high latency in current location

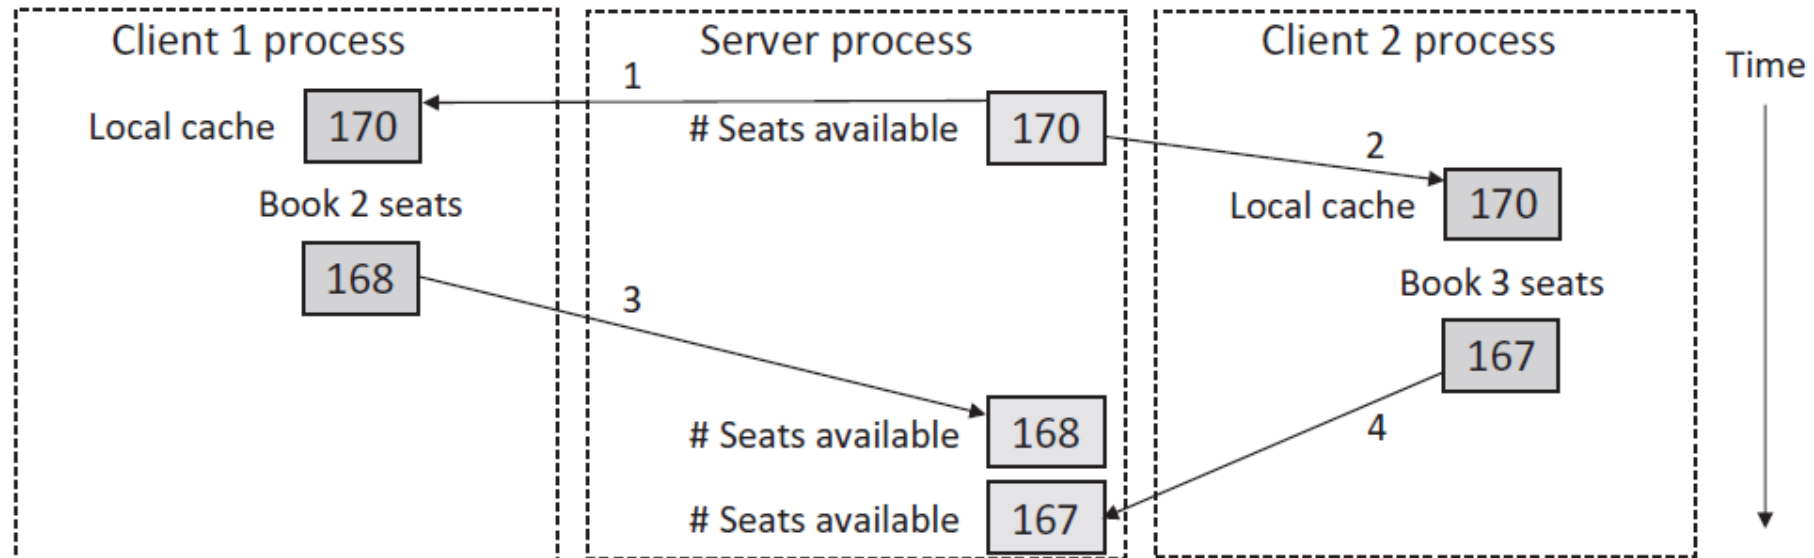
# Replication Transparency

- Requires that **multiple copies of objects** can be created without any effect of the replication **seen by applications** that uses these objects
- The most significant challenge – maintenance of **consistency**



# Concurrency Transparency

- Requires **that concurrent processes** can share objects without interference
- Raises the issue of data consistency
- Example of lost update problem – an airline seat booking scenario:



# Failure Transparency

- Failure transparency requires that **faults are concealed** so that applications can **continue to function** without any impact on behavior or correctness
- Failures in distributed systems are inevitable
- We need to rely on **runtime techniques** that will deal with the failures
- Example – TCP and UDP
  - TCP has a **number of built-in features** that transparently deal with various failure
  - UDP is a more lightweight protocol – “send and pray”

# Degree of Transparency

Aiming at full distribution transparency may be too much:

- Users may be located in **different continents**
- **Completely hiding failures** of networks and nodes is (theoretically and practically) impossible
  - You cannot distinguish a slow computer from a failing one
  - You can never be sure that a server actually performed an operation before a crash
- Full transparency will **cost performance**, exposing distribution of the system
  - Keeping Web caches exactly up-to-date with the master
  - Immediately flushing write operations to disk for fault tolerance

# Open Distributed Systems

- Offer services that are described a priori
  - Syntax and semantics are known via protocols
- Services specified via interfaces
- Benefits
  - Interoperability
  - Portability
  - Extensibility
    - Open system evolve over time and should be extensible to accommodate new functionality.

# Scalability

Many developers of distributed system easily use the adjective “scalable” without making clear **why** their system actually scales.

## Dimensions of scalability:

- Number of users and/or processes (**size scalability**)
- Maximum distance between nodes (**geographical scalability**)
- Number of administrative domains (**administrative scalability**)

Most systems account only, to a certain extent, for size scalability.



# Scalability (cont.)

## **Limitation of size scalability**

- Centralized services
- Centralized data
- Centralized algorithms

# Scalability (cont.)

## **Limitations of geographical scalability**

- Synchronous communication
- Networks is unreliable
- Centralized components in the system

# Scalability (cont.)

## **Limitations of administrative scalability**

- Conflicting policies
  - Resource usage
  - Management
  - Security

# Scaling Techniques

**Hide communication latencies** – avoid waiting for responses; do something else:

- Make use of asynchronous communication
- Problem: not every application fits this model

# Scaling Techniques (cont.)

**Distribution** – partition data and computations across multiple machines:

- Move computations to clients
- Decentralized naming services
- Decentralized information systems

# Scaling Techniques (cont.)

**Replication/caching** – make copies of data available at different machines:

- Replicated file servers and databases
- Mirrored Web sites
- Web caches (in browsers)
- File caching (at server and client)

# Scaling – Challenges

Applying scaling techniques seems easy, except for:

- Having multiple copies (cached or replicated), leads to **inconsistencies**: modifying one copy makes that copy different from the rest
- Always keeping copies consistent and in a general way requires **global synchronization** on each modification

If we can tolerate inconsistencies, we may reduce the need for global synchronization, but **tolerating inconsistencies is application dependent**.

# Summary: Scaling Techniques

- *Principles* for good decentralized algorithms
  - No machine has a complete state
  - Make decisions based on local information
  - A single failure does not bring down the system
  - No global clock
- *Techniques*
  - Asynchronous communication
  - Distribution
  - Caching and replication



# Challenges and Pitfalls

Many distributed systems are needlessly complex due to **mistakes** that required patching later on. There are many **false assumptions**:

- The network is reliable
- The network is secure
- The network is homogeneous
- The topology does not change
- Latency is zero
- Bandwidth is infinite
- Transport cost is zero
- There is one administrator

# References

- Distributed Systems: Principles and Paradigms by Tanenbaum and van Steen, chapter 1