

# Lecture#2: Distributed Systems Models and Architectures

S. M. Ahsan Kazmi

# Outline

- Traditional DS Models
- Emerging DS Models
- Uniprocessor and Distributed OS
- Architectural Styles
- System Architecture

# Distributed Systems Models

- **Minicomputer model** (e.g., early networks)
  - Each user has a local machine
  - Local processing but can fetch remote data (files, databases)
- **Workstation model** (e.g., Sprite)
  - Processing can also migrate
- **Client-server Model** (e.g., V system, world wide web)
  - User has a local workstation
  - Powerful workstations serve as servers (file, print, DB servers)

# Distributed System Models (cont.)

- **Cluster computing systems / Data centers**
  - LAN with a cluster of servers + storage
    - Linux, Mosix, ..
    - Used by distributed web servers, scientific applications, enterprise applications
- **Grid computing systems**
  - Cluster of machines connected over a WAN
  - SETI @ home
- **WAN-based clusters / distributed data centers**
  - Google, Amazon, ...

# Emerging Models

- **Distributed Pervasive Systems**
  - “smaller” nodes with networking capabilities
    - Computing is “everywhere”
  - Home networks: Windows Media Center, ...
  - Mobile computing: smartphones, iPods, Car-based PCs
  - Sensor networks


# Distributed Pervasive Systems

Overlapping subtypes:

- **Ubiquitous computing systems:**
  - pervasive and continuously present, i.e., there is a continuous interaction between system and user.
- **Mobile computing systems:**
  - pervasive, but the emphasis is on the fact that devices are inherently mobile.
- **Sensor (and actuator) networks:**
  - pervasive, with emphasis on the actual (collaborative) sensing and actuation of the environment.

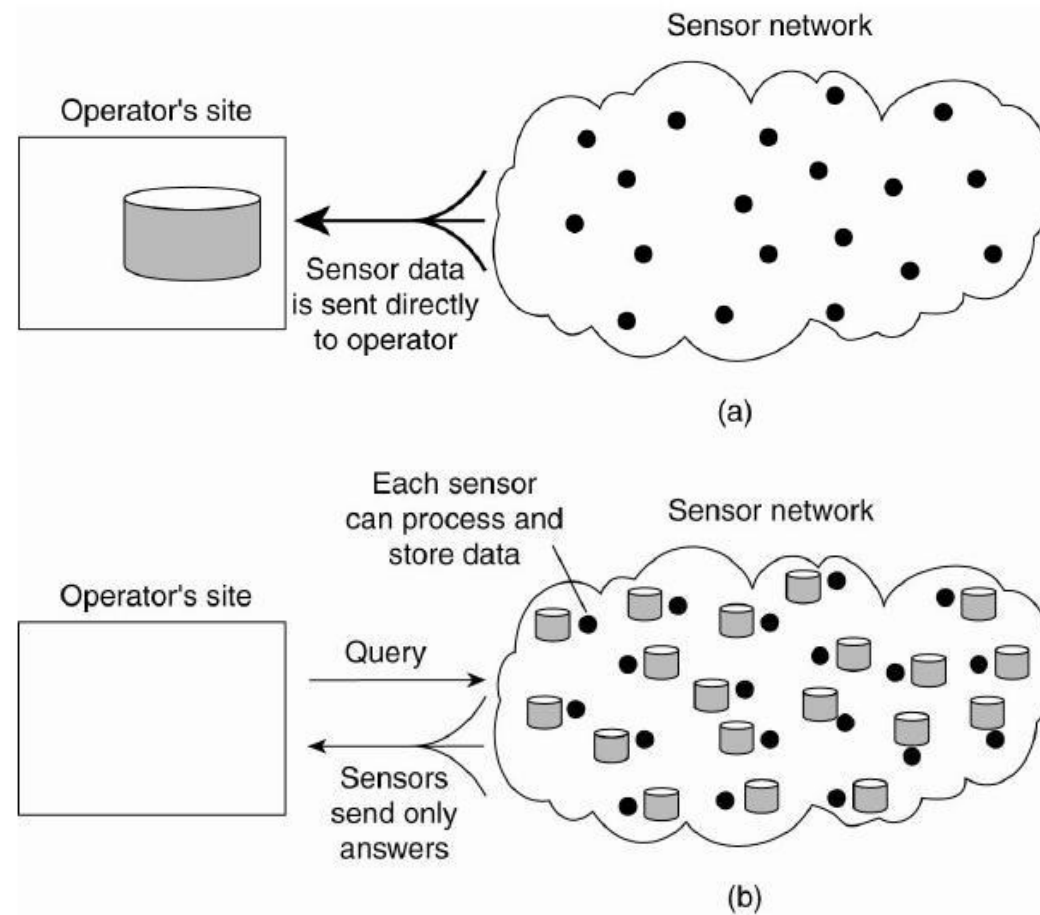
# Distributed Pervasive systems (cont.)

## Ubiquitous computing systems



# Distributed Pervasive Systems (cont.)

## Sensor network





# Uniprocessor Operating Systems

- An OS acts as a **resource manager** or an arbitrator
  - Manages CPU, I/O devices, memory
- OS provides a **virtual interface** that is easier to use than hardware
- Structure of uniprocessor operating systems
  - Monolithic design (e.g., MS-DOS, early UNIX)
    - One large kernel that handles everything
  - Layered design
    - Functionality is decomposed into N layers
    - Each layer uses services of layer N-1 and implements new service(s) for layer N+1

# Multiprocessor Operating Systems

- Like a uniprocessor operating system
- Manages **multiple CPUs transparently** to the user
- Each processor has its own hardware cache
  - Maintain consistency of cached data

# Distributed Operating System

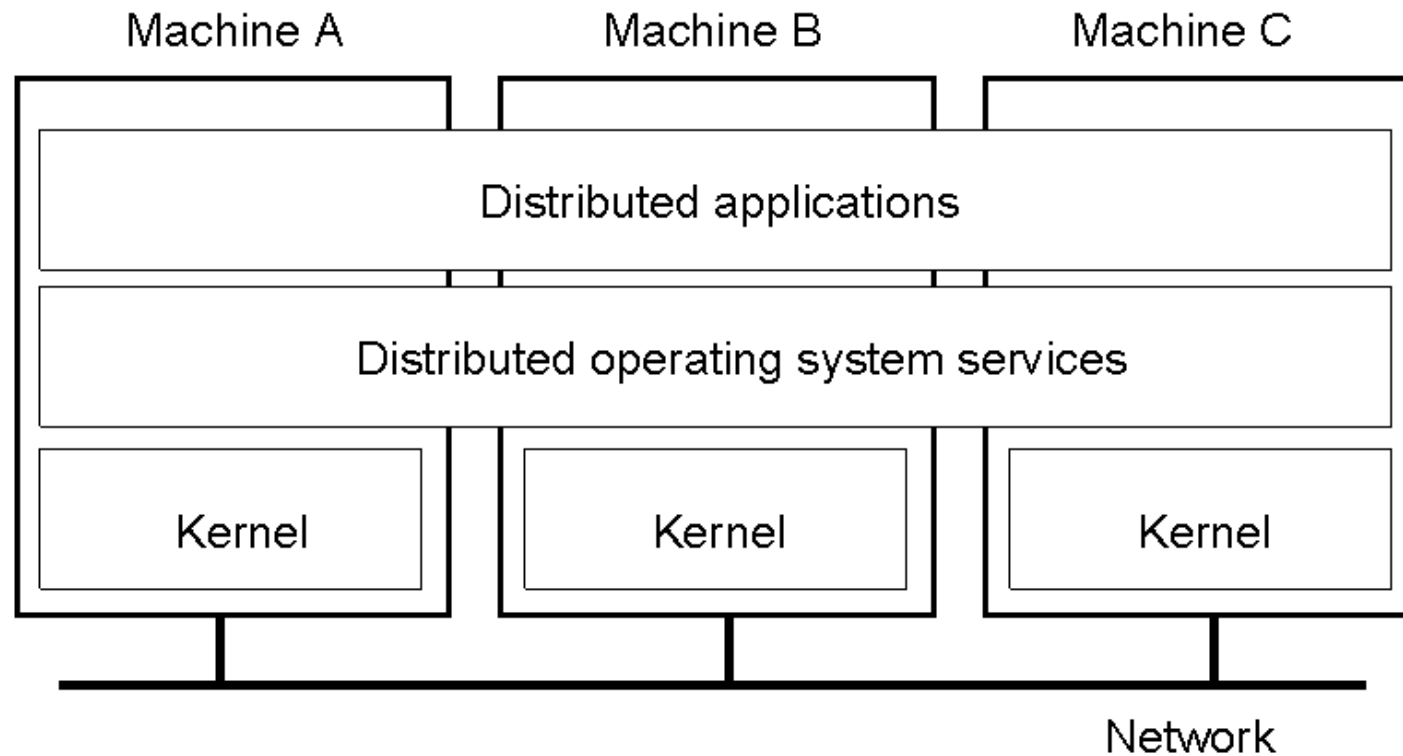
- Manages resources in a distributed system
  - Seamlessly and transparently to the user
- Looks to the user like a centralized OS
  - But operates on multiple independent CPUs
- Provides transparency
  - Location, migration, concurrency, replication,...
- Presents users with a virtual uniprocessor

# Types of Distributed OSs

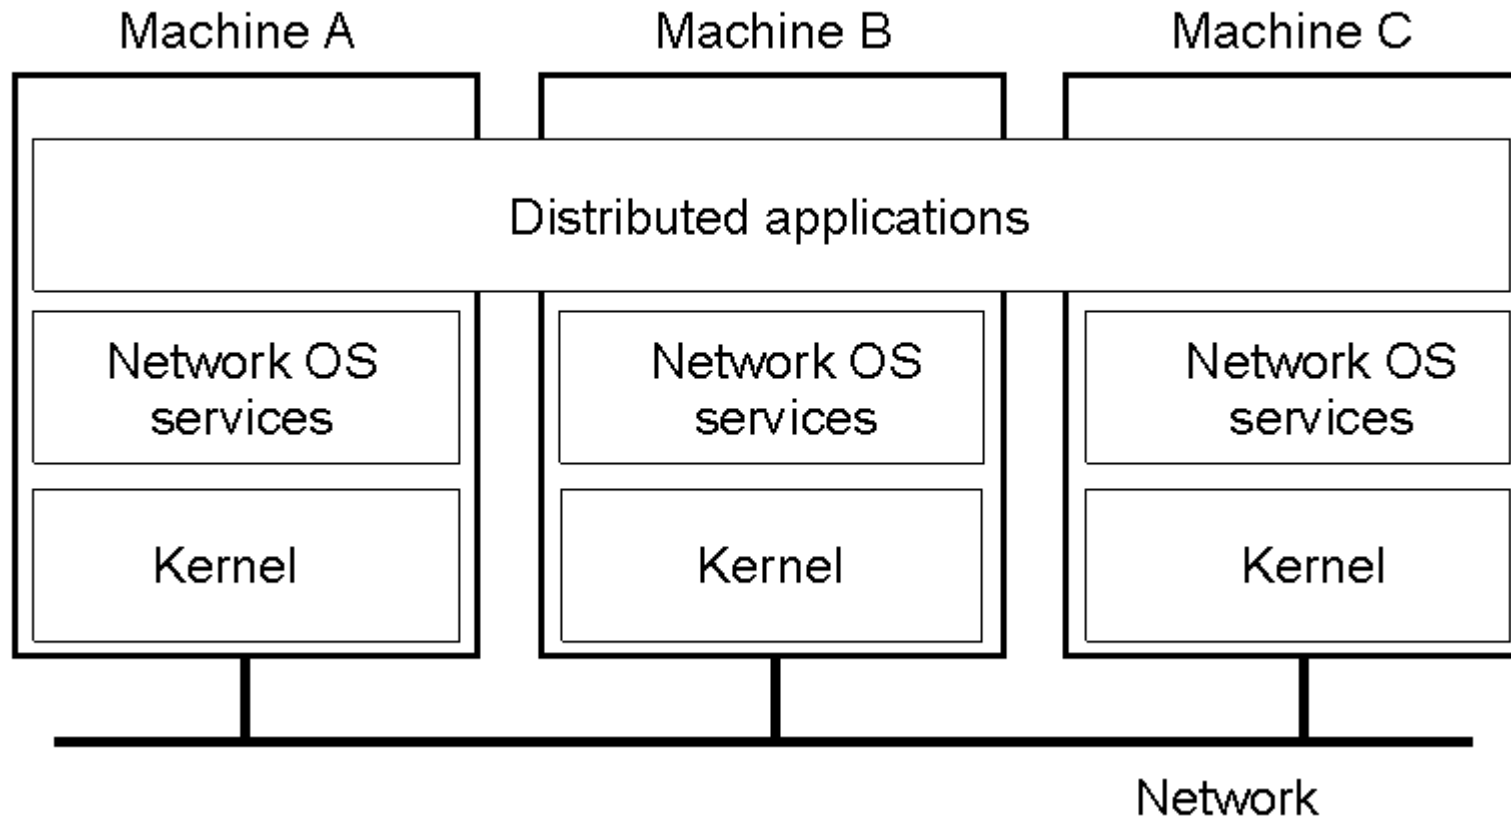
System	Description	Main Goal
DOS	A tightly-coupled operating system for multi-processors and homogeneous multicomputers	Hide and manage hardware resources
NOS	Loosely-coupled operating system for heterogeneous multicomputers (LAN and WAN)	Offer local services to remote clients
Middleware	Additional layer on top of NOS implementing general-purpose services	Provide distribution transparency

# Multicomputer Operating Systems

Example: MOSIX cluster - single system image



# Network Operating System



# Comparison between Systems

Item	Distributed OS		Network OS
	Multiproc.	Multicomp.	
Degree of transparency	Very High	High	Low
Same OS on all nodes	Yes	Yes	No
Number of copies of OS	1	N	N
Basis for communication	Shared memory	Messages	Files
Resource management	Global, central	Global, distributed	Per node
Scalability	No	Moderately	Yes
Openness	Closed	Closed	Open


# Architectural Styles

- Architectural style describes a particular way **how we can organize a collection** of components distributed over the various machines
- Important styles of architecture for distributed systems
  - Layered
  - Object-based
  - Event-based
  - Resource-centered




# Layered Architectures

Layered system organizations: Each layer uses **previous layer** to implement new functionality that is exported to the **layer above**



# Object-based Architectures


- Each object correspond to a component
- Typically, each object contains
  - Set of methods
  - State(s)
  - Interface



# Event-based Architecture

Communicate via a common repository

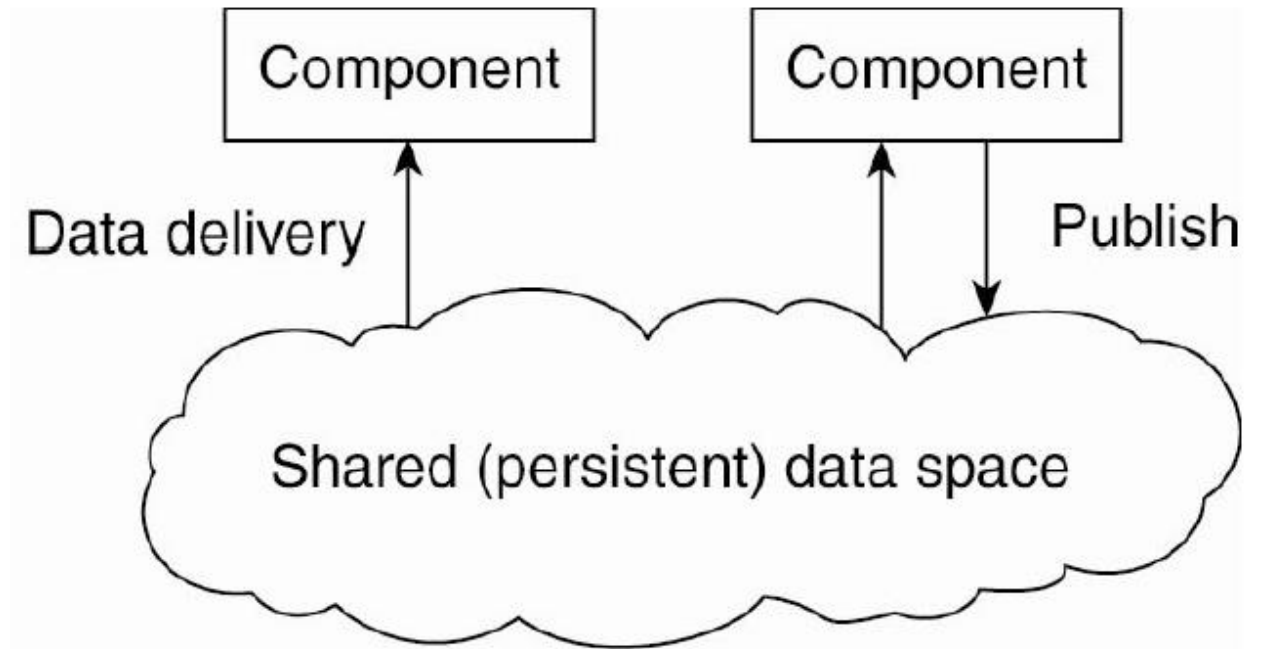
- Use a publish-subscribe paradigm
- Consumers subscribe to types of events
- Events are delivered once published by any publisher



# Shared Data-space Architectures

“Bulletin-board” architecture

- Decoupled in space and time
- Post items to shared space; consumers pick up at a later time



# System Architecture

**System architecture** – describe the placement of software components on physical machines

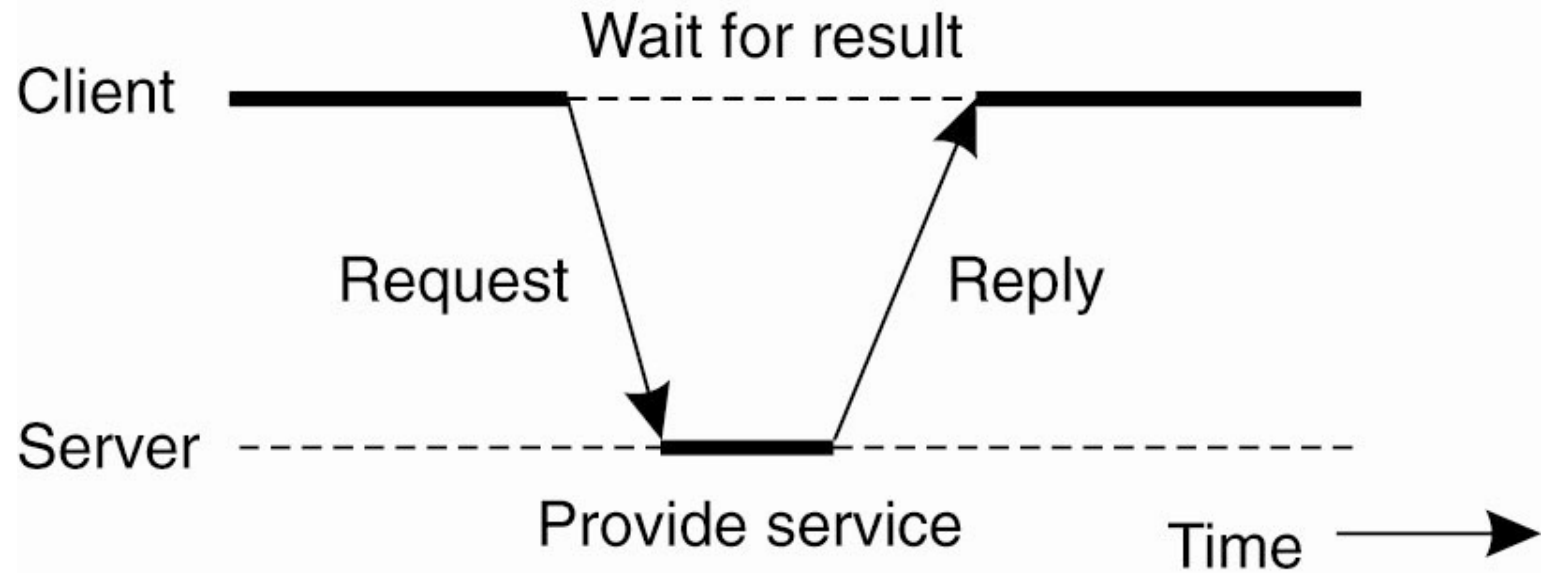
The realization of architecture may be

- **Centralized.** Client-server architecture
  - **Decentralized.** Peer-to-peer architecture
  - **Hybrid.** Combination of Centralized and Decentralized architectures
- (according to Tanenbaum and van Steen)

# Centralized Organizations

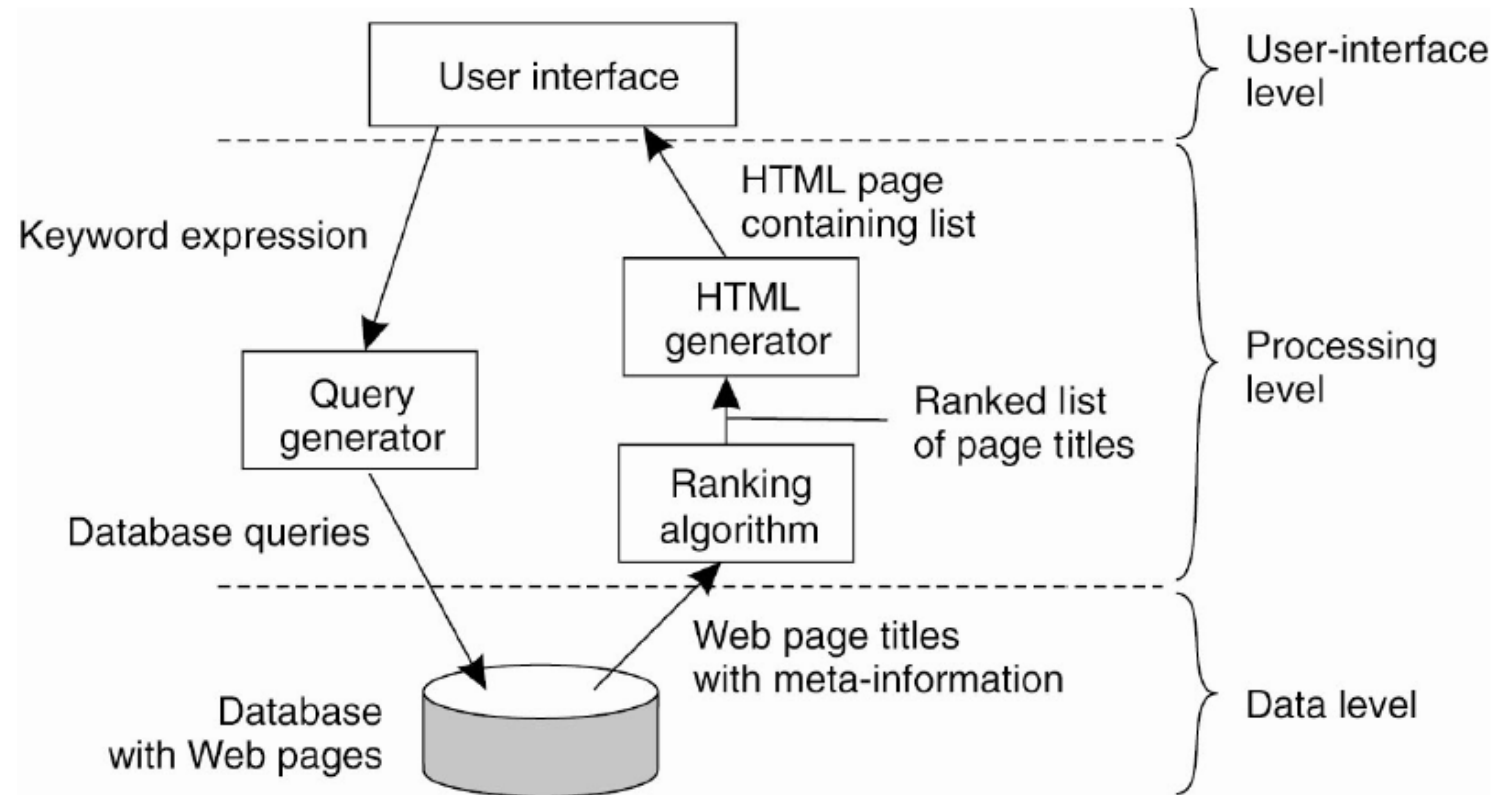
Traditional client-server architecture

- Logical separation on server and client functionality
- Application layering
  - User-interface level
  - Processing level
  - Data level



# Application Layering

- **User-interface level** contains units for an application's UI
- **Processing level** contains the functions of an application
- **Data level** contains the data that user wants to manipulate through the application components



Search engine architecture with 3 layers

# Multi-tiered Architectures

- The simplest organization is to have only two types of machines:
- A client machine containing only the programs implementing (part of) the user interface level
- A server machine containing the rest,
  - the programs implementing the processing and data level

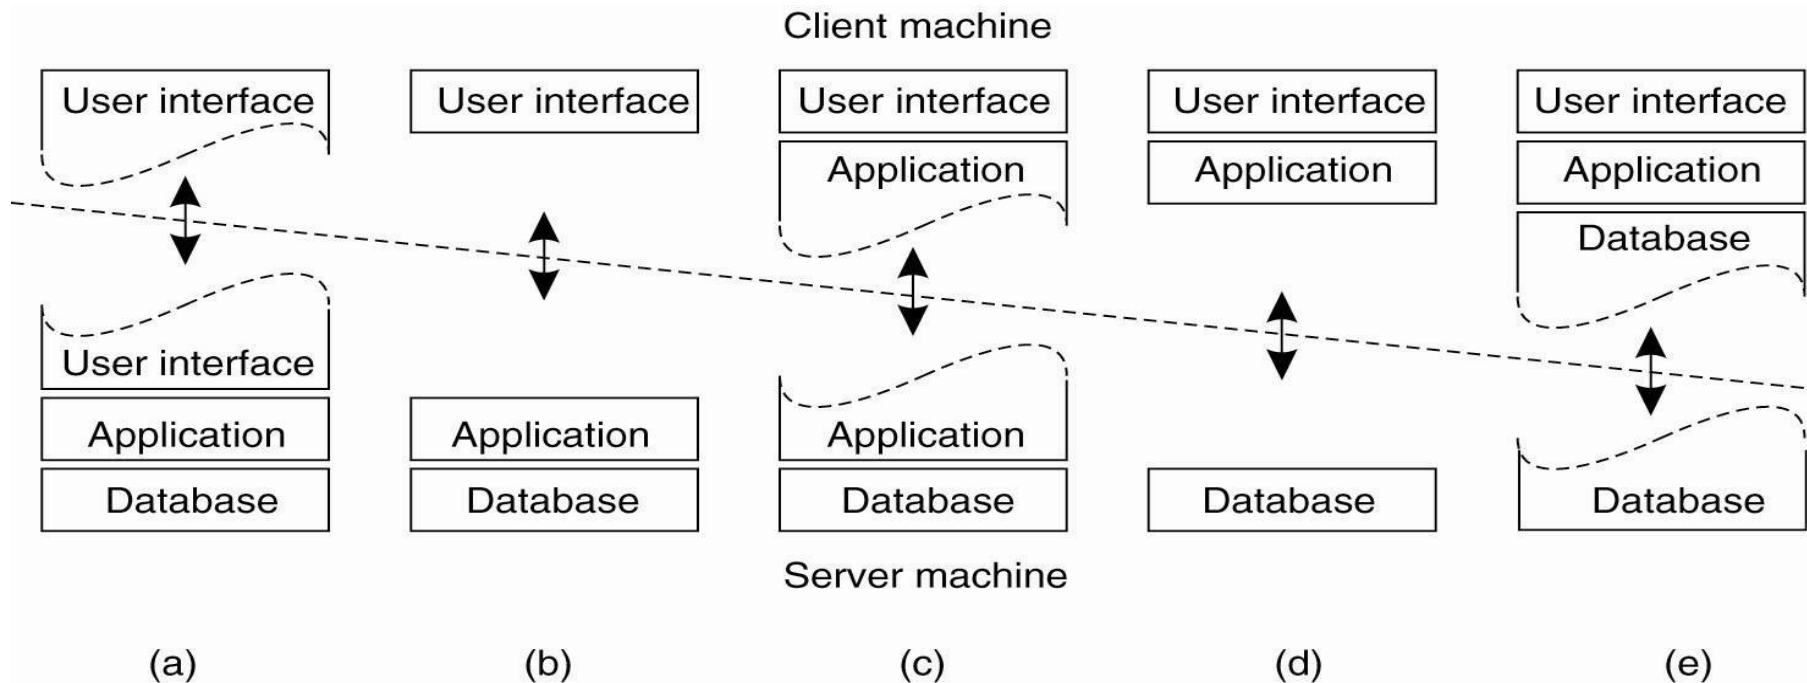


# Two-tiered: Thin vs Fat client

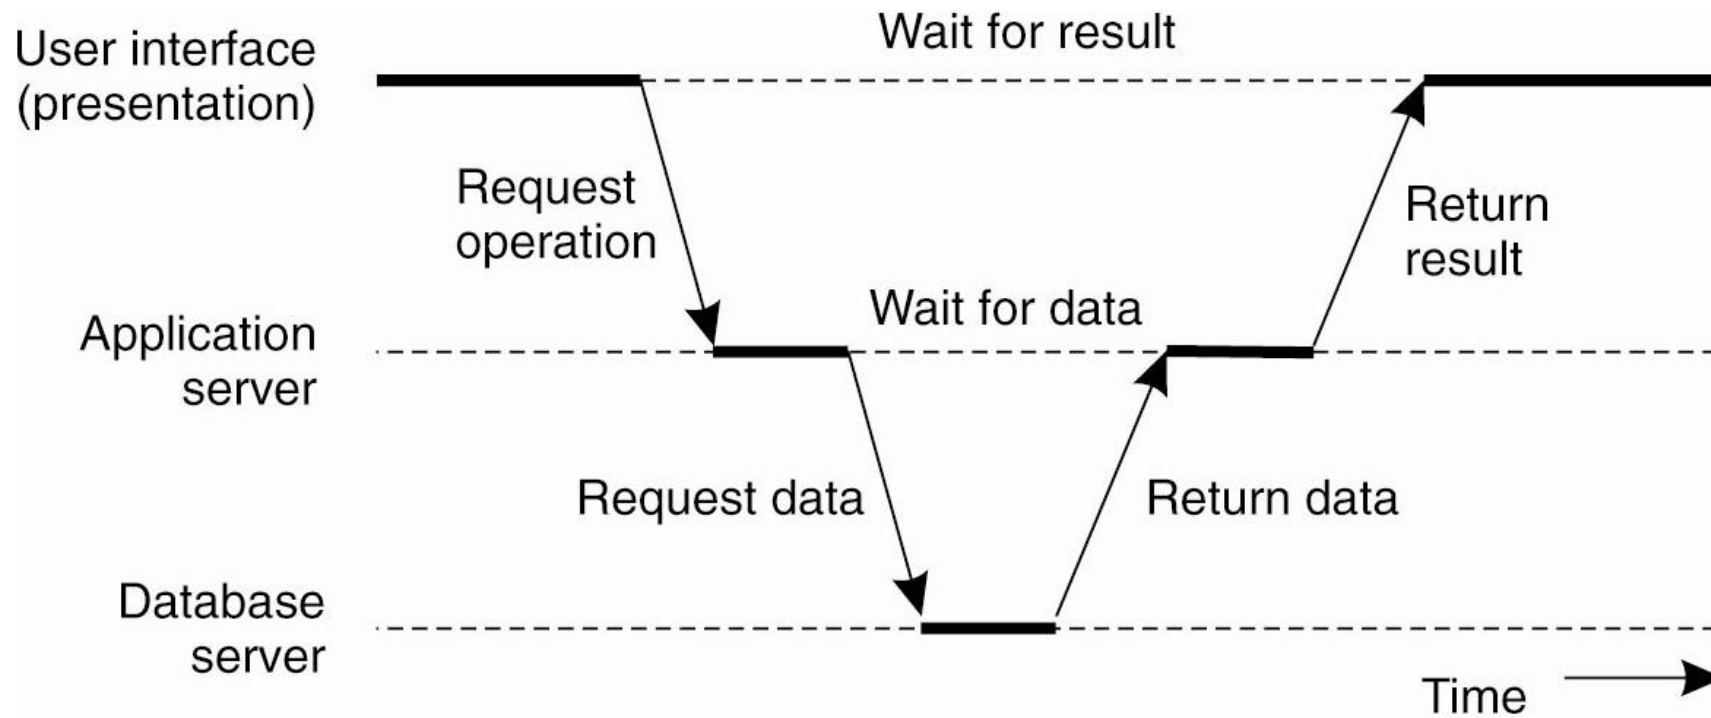
- Thin client
  - Server provides processing and data management
  - Client provides a simple graphical display
  - No performance issues with the client
  - Easier to manage, more reliable, less money for the helpdesk
- Fat client
  - Some data processing on the client
  - Reduces workload at a server; more scalable
  - Harder to manage by a system administrator

# A Spectrum of Choices

- Single-tiered – mainframe configuration
- Two-tiered – client/single server configuration
- Three-tiered – several server configurations



# Three-tier Web Applications



# Decentralized Architectures


- **Structured P2P:** nodes are organized as a specific distributed structure
- **Unstructured P2P:** nodes randomly select neighbors
- **Hierarchical P2P:** some nodes have special functions

The overlay network connects nodes in the P2P system.

- Nodes in the overlay use their own addressing system for data handling
- Nodes may forward requests to locations that may not be known by the requester

# Overlay Network

- Logical network which built on top of a physical network
- Each node in the P2P system knows how to contact several other nodes
- The overlay network may be
  - structured (overlay network constructed using a deterministic procedure)
  - unstructured (randomized algorithms)



# P2P System Characteristics

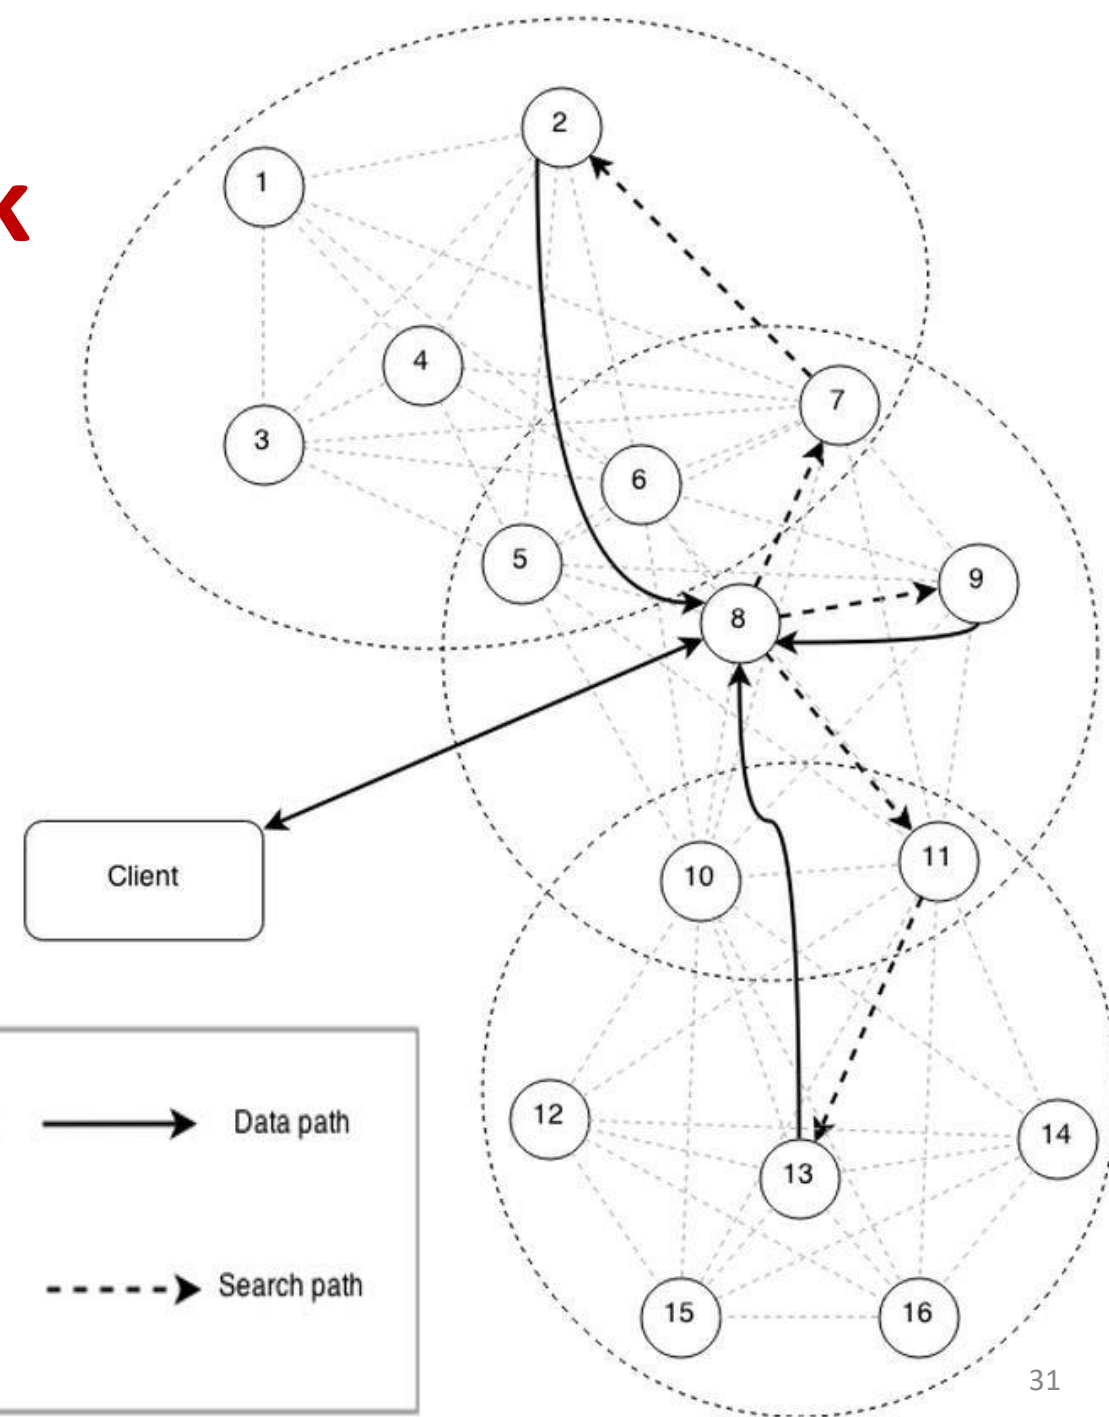
- The design ensures that each user contributes resources to the system
- All the nodes in a peer-to-peer system have the **same functional capabilities** and responsibilities
- The correct operation does not depend on the existence of any centrally administered systems (pure p2p)
- Key issues & challenges
  - Choice of an algorithm for the placement of data across many hosts
  - Balance the workload and ensures availability without adding extra overhead

# Schema of a P2P Network

Nodes can:

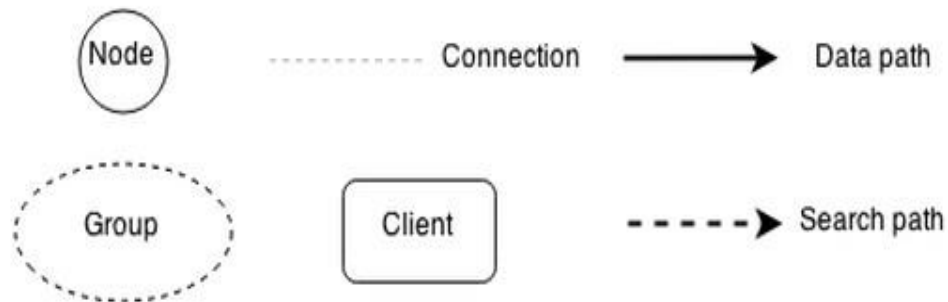
- join the system
- leave the system
- publish a resource
- search a resource
- unpublish a resource

Resources are replicated



## Node Groups

This view shows organization of single storage node into groups, groups interconnection, data search paths and data flow paths, principles of client connection to a storage node.



# Structured P2P Network

- **Basic features** of structured P2P overlays:
  - Structure – to accommodate participating nodes and data in the overlay
    - Structured overlays use several different geometries (rings, trees, hypercubes, etc.)
  - Routing algorithm – to locate nodes in the overlay and insert/retrieve data to/from them
  - Join/leave mechanisms – to enable self-organization and fault tolerance
  - The **primary goal is to enable the deterministic lookup** (i.e., access guarantees with certain time bounds)



# Example: Chord

- Chord was developed at MIT
- Originally published in 2001 at the Sigcomm conference
- Chord's overlay routing principle quite easy to understand
  - Paper has mathematical proofs of correctness and performance
- Uses DHT to locate objects
- Many projects at MIT around Chord

# Unstructured P2P

- Unstructured P2P organizes the overlay network by using randomized Algorithms
- Each node knows about a subset of nodes
  - The subset could be chosen in very different ways: physically close nodes, nodes that joined at about the same time, etc.
- Data randomly mapped to some nodes in the system

# Unstructured P2P: Locating the data

- **Flooding.** Node A sends a lookup query to all its neighbors. A neighbor responds, or forwards (floods) the request.
  - Limited flooding (maximum number of forwarding) TTL
  - Probabilistic flooding (flood only with a certain probability)
- **Random walk.** Randomly select a neighbor B. If B has the answer, it replies, otherwise B randomly selects one of its neighbors.
  - Parallel random walk - works well with replicated data


# Structured vs Unstructured

- Structured networks typically guarantee that if an object is in the network it will be available in a bounded amount of time
- Unstructured networks offer no guarantees.
  - For example, some will only forward search requests a specific number of hops
  - Random graph approach means there may be loops
  - Graph may become disconnected

# Hierarchically organized P2P Networks

Superpeers (in some sources ultrapeers)


- Maintain indexes to some or all nodes in the system
  - Supports resource discovery
  - Act as servers to regular peer nodes, peers to other superpeers
- Improve scalability by controlling floods
- Can also monitor the state of the network
- Example: Skype



# Hybrid Architectures

- Combine client-server and P2P architectures
  - **Edge-server systems**; e.g., ISP, which act as servers to their clients, but cooperate with other edge servers to host shared content
  - **Collaborative distributed systems**; e.g., BitTorrent, which supports parallel downloading and uploading chunks of data.
    - First, interact with client-server system, then operate in decentralized manner

# Edge-server Systems



Viewing the Internet as consisting of a collection of edge servers

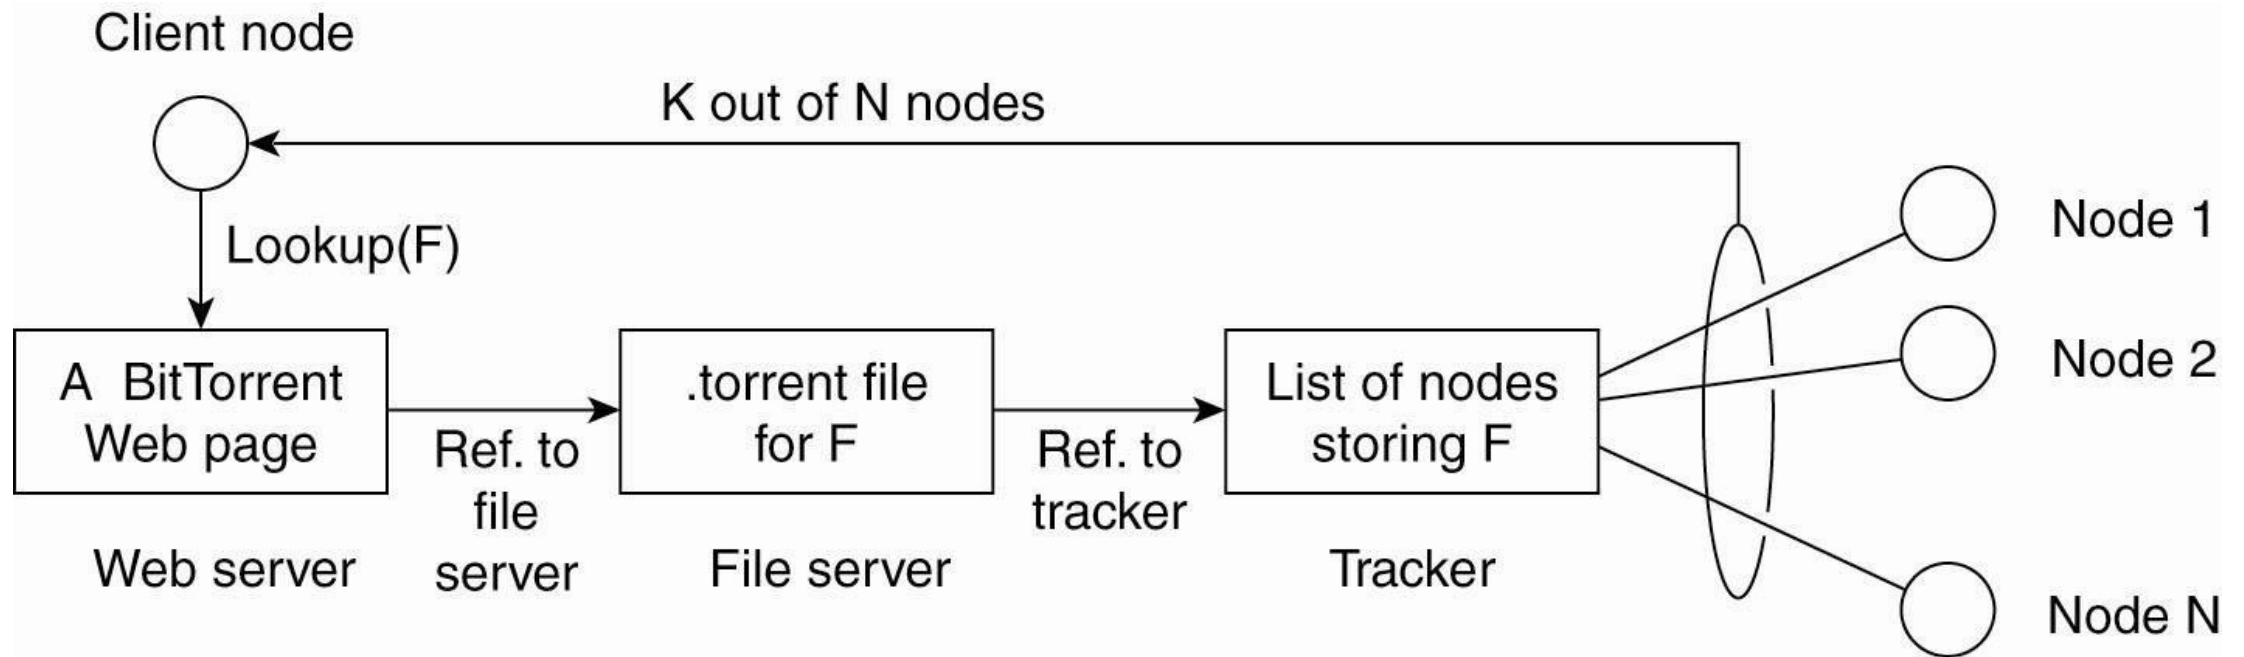
# Collaborative Distributed Systems

Example: BitTorrent

- Clients contacts a global directory (Web server) to locate
  - a *.torrent* file with the information needed to locate a **tracker**
  - a server that can supply a list of **active nodes (peers)** that have chunks of the desired file
- Using this information, clients can download the file in chunks from multiple sites in the network
- Clients must also provide file chunks to other users



# BitTorrent



# References

- Distributed Systems: Principles and Paradigms by Tanenbaum and van Steen, chapter 1
- Distributed Systems: Principles and Paradigms by Tanenbaum and van Steen, chapter 2

# Thank you!

Questions?