

Lecture#6 –Communication in DS

S. M. Ahsan Kazmi

Outline

- Communication in DS
 - Introduction
 - Types of communication
 - Communication Models
 - Remote procedure calls

Recap

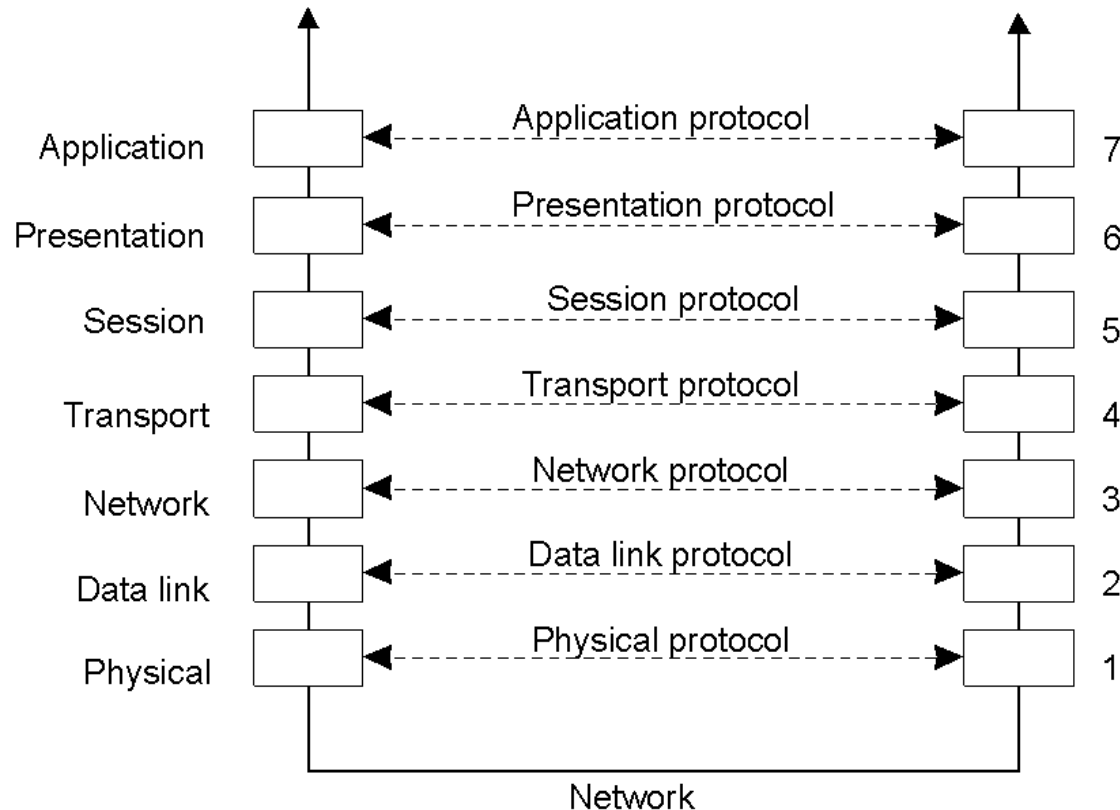
- OS-Level Virtualization
 - Comparison with different types of virtualization
 - Containers Building Blocks
 - Containers Orchestration
- Memory Technologies and Hierarchies
 - Virtualization- Memory Basics and Challenges
 - Memory reclamation approaches
- Data Center and Cloud
 - Inside a data center
 - Data center economics
 - Challenges
 - Service and Deployment models

Communication Between Processes

- Unstructured communication
 - Use shared memory or shared data structures
- Structured communication
 - Use explicit messages (IPCs)
- Distributed Systems: both need low-level communication support

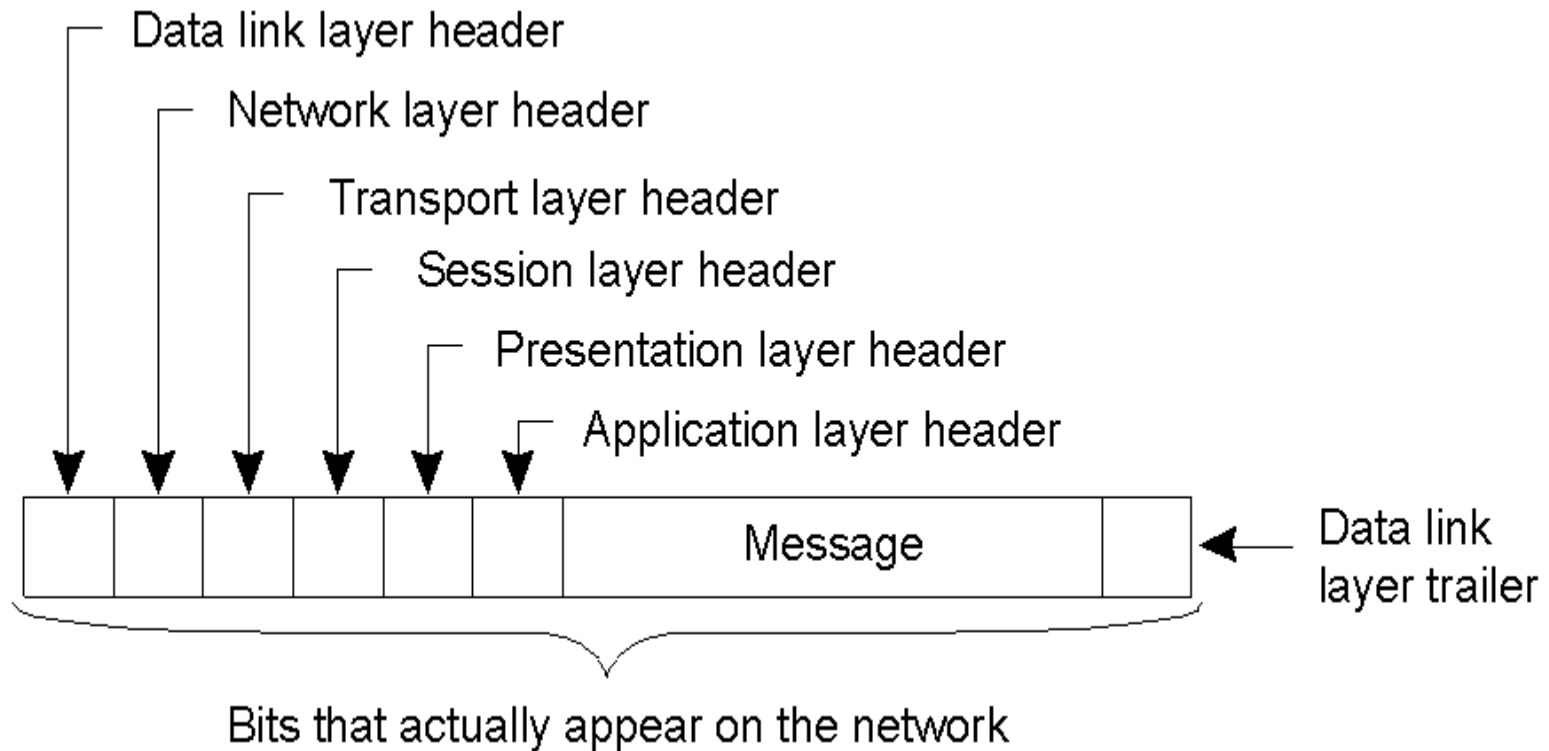
Layered Protocols

- **OSI**: Open Systems Interconnection
- Developed by the International Organization for Standardization (ISO)
- Provides a **generic framework** to discuss the layers and functionalities of communication protocols.



Layers, interfaces, and protocols in the OSI model.

OSI Model (con.t)



A typical message as it appears on the network.

OSI Protocol Layers

- **Physical layer**

- Deals with the **transmission of bits**
- Physical interface between data transmission device
- (e.g. computer) and transmission medium or network
- **Concerned with:**
 - Characteristics of transmission medium, Signal levels, Data rates

- **Data link layer:**

- Deals with detecting and correcting bit transmission errors
- Bits are group into **frames**
- **Checksums** are used for integrity

OSI Protocol Layers (con.t)

- **Network layer:**

- Performs multi-hop **routing** across multiple networks
- Implemented in end systems and routers

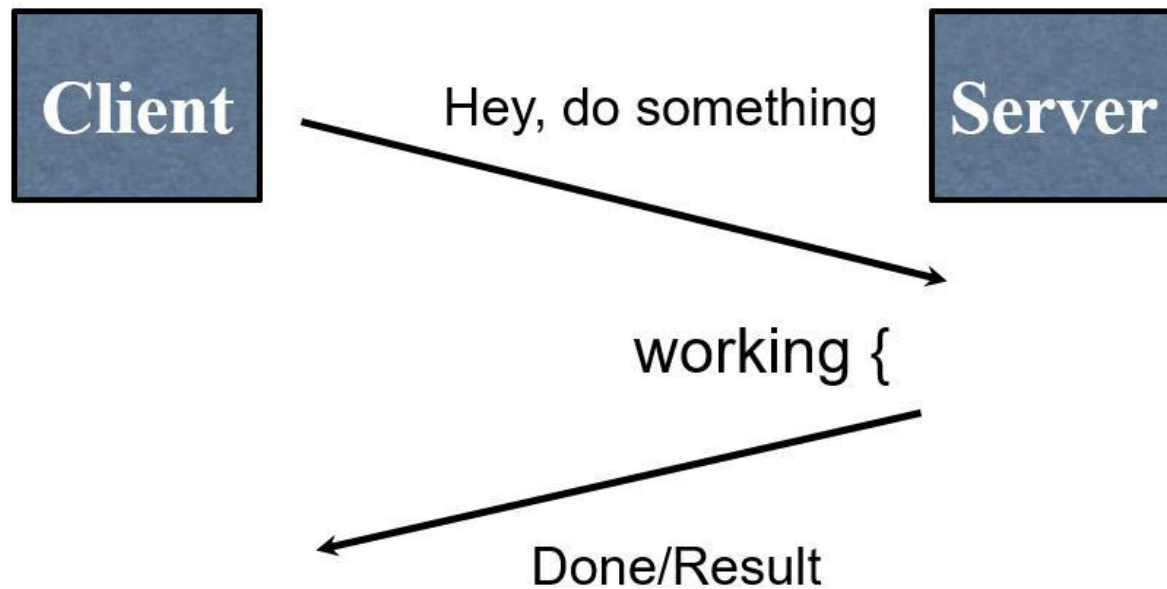
- **Transport layer:**

- Packing of data
- Reliable delivery of data (breaks message into pieces small enough, assign each one a sequence number and then send them)
- Ordering of delivery
- **Examples:**
 - TCP (connection-oriented)
 - UDP (connectionless)

OSI Protocol Layers (con.t)

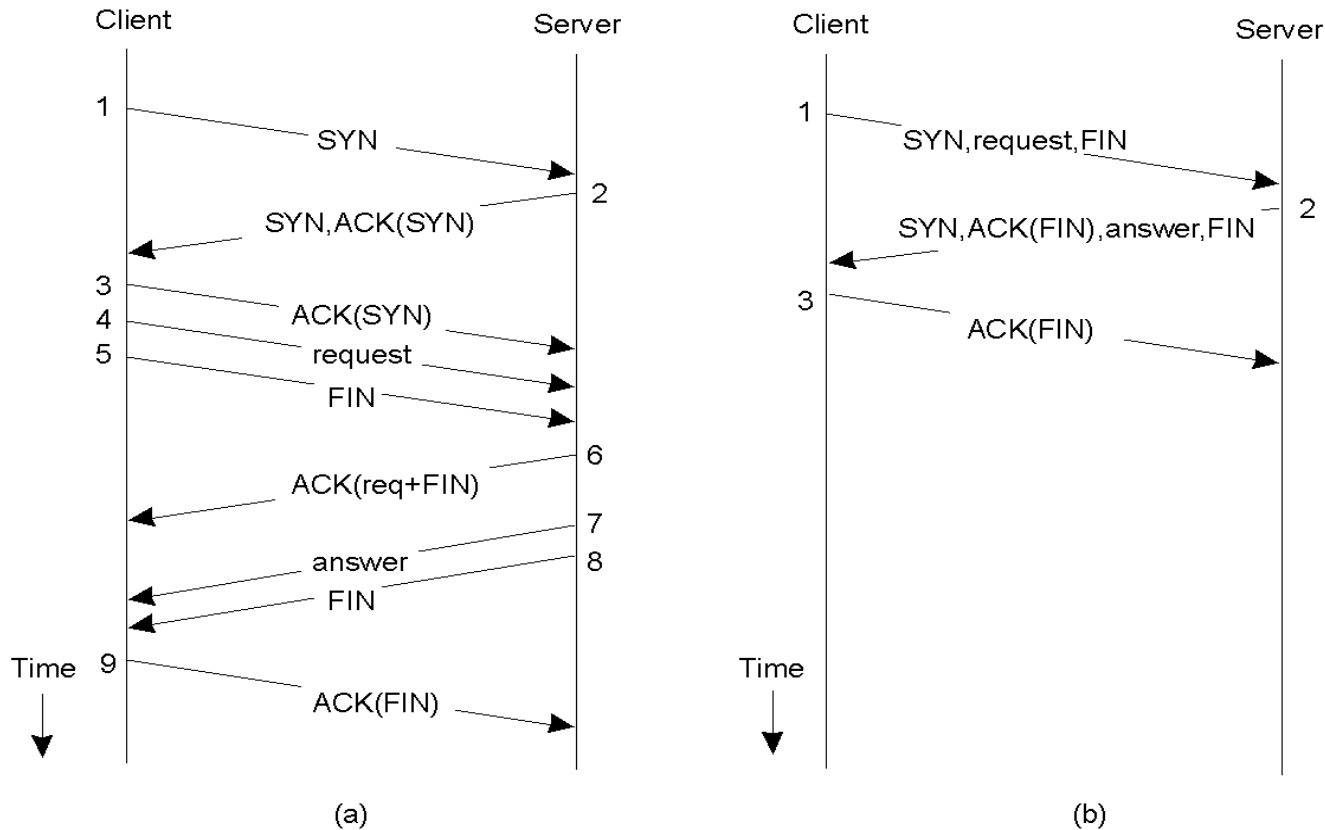
- **Session layer**
 - Provide **dialog control** to keep track of which party is talking and it provides synchronization facilities
 -
- **Presentation layer**
 - Deals with non-uniform **data representation** and with **compression** and **encryption**
- **Application layer**
 - Support for user applications
 - e.g. HTTP, SMTP, FTP

Typical Communication Pattern



OSI Protocol Layers (con.t)

Client-Server TCP protocol



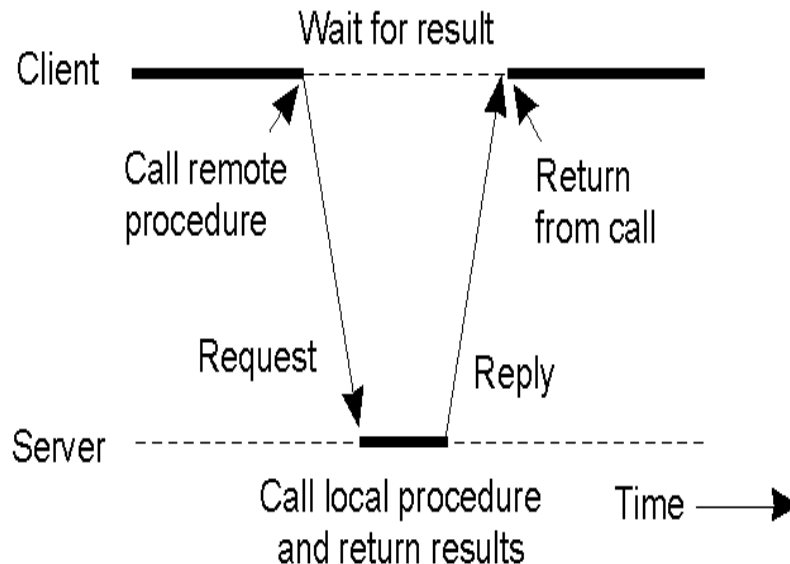
(a) Normal operation of TCP. (b) Transactional TCP.

Communication Models: Push or Pull ?

- Client-pull architecture
 - Clients pull data from servers (by sending requests)
 - Example: HTTP
- Server-push architecture
 - Servers push data to client
 - Example: video streaming

Remote Procedure call

- **Basic idea:** To execute a **procedure at a remote site** and ship the results back.
- **Goal:** To make this operation as distribution transparent as possible (i.e., the remote procedure call should look like a local one to the calling procedure).



RPC Goals

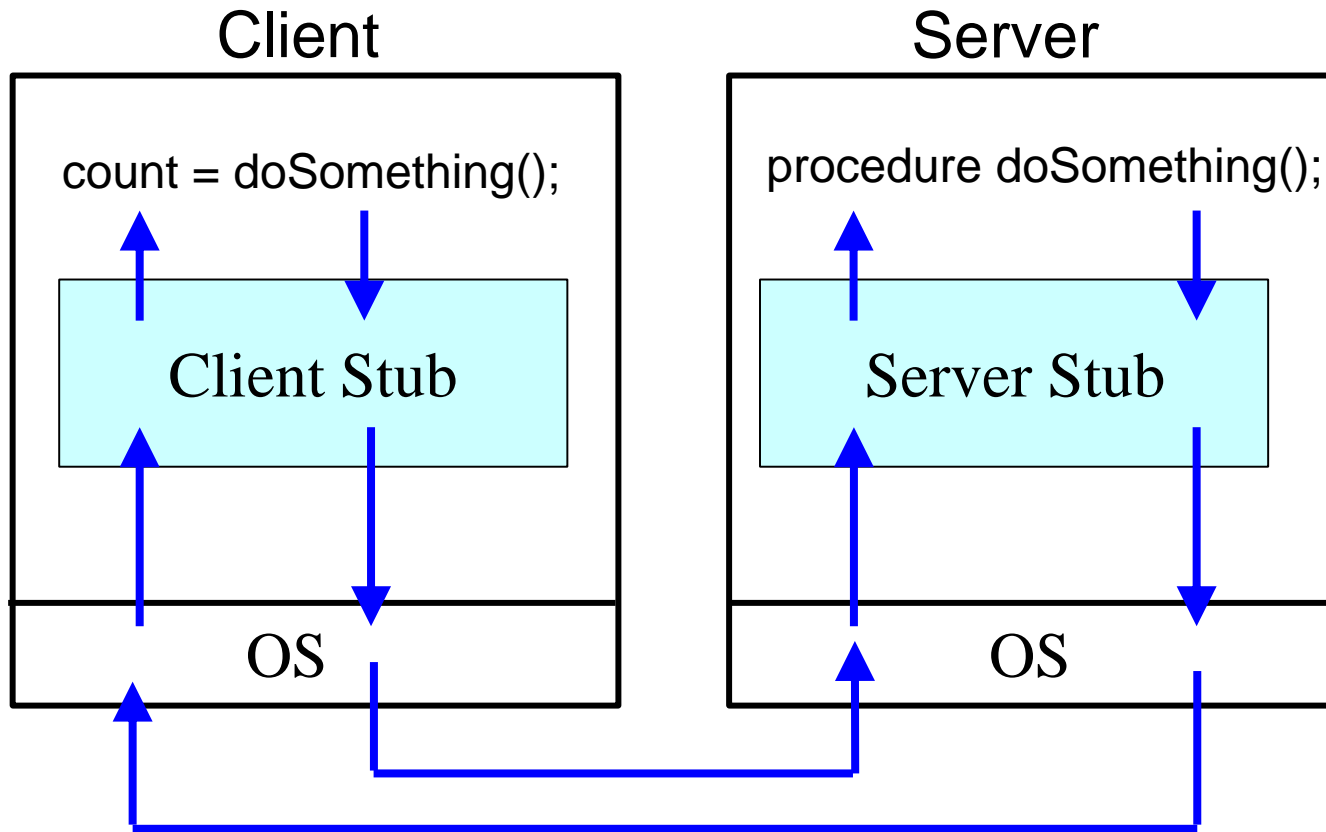
- A remote procedure call makes a call to a remote service look like a local call
 - RPC makes transparent whether the server is local or remote
 - RPC makes the architecture of remote machines transparent
- Benefits
 - Ease of programming
 - Hide complexity
 - Automates task of implementing distributed computation
 - Familiar model for programmers (just make a function call)

Parameter Passing

- Local procedure parameter passing
 - Call-by-value
 - Call-by-reference: arrays, complex data structures
- Remote procedure calls simulate this through:
 - Stubs
 - Marshalling
- Related issue: global variables are not allowed in RPCs

Client and Server Stubs

Stubs are additional functions that are added to the main functions in order to **support for RPC**



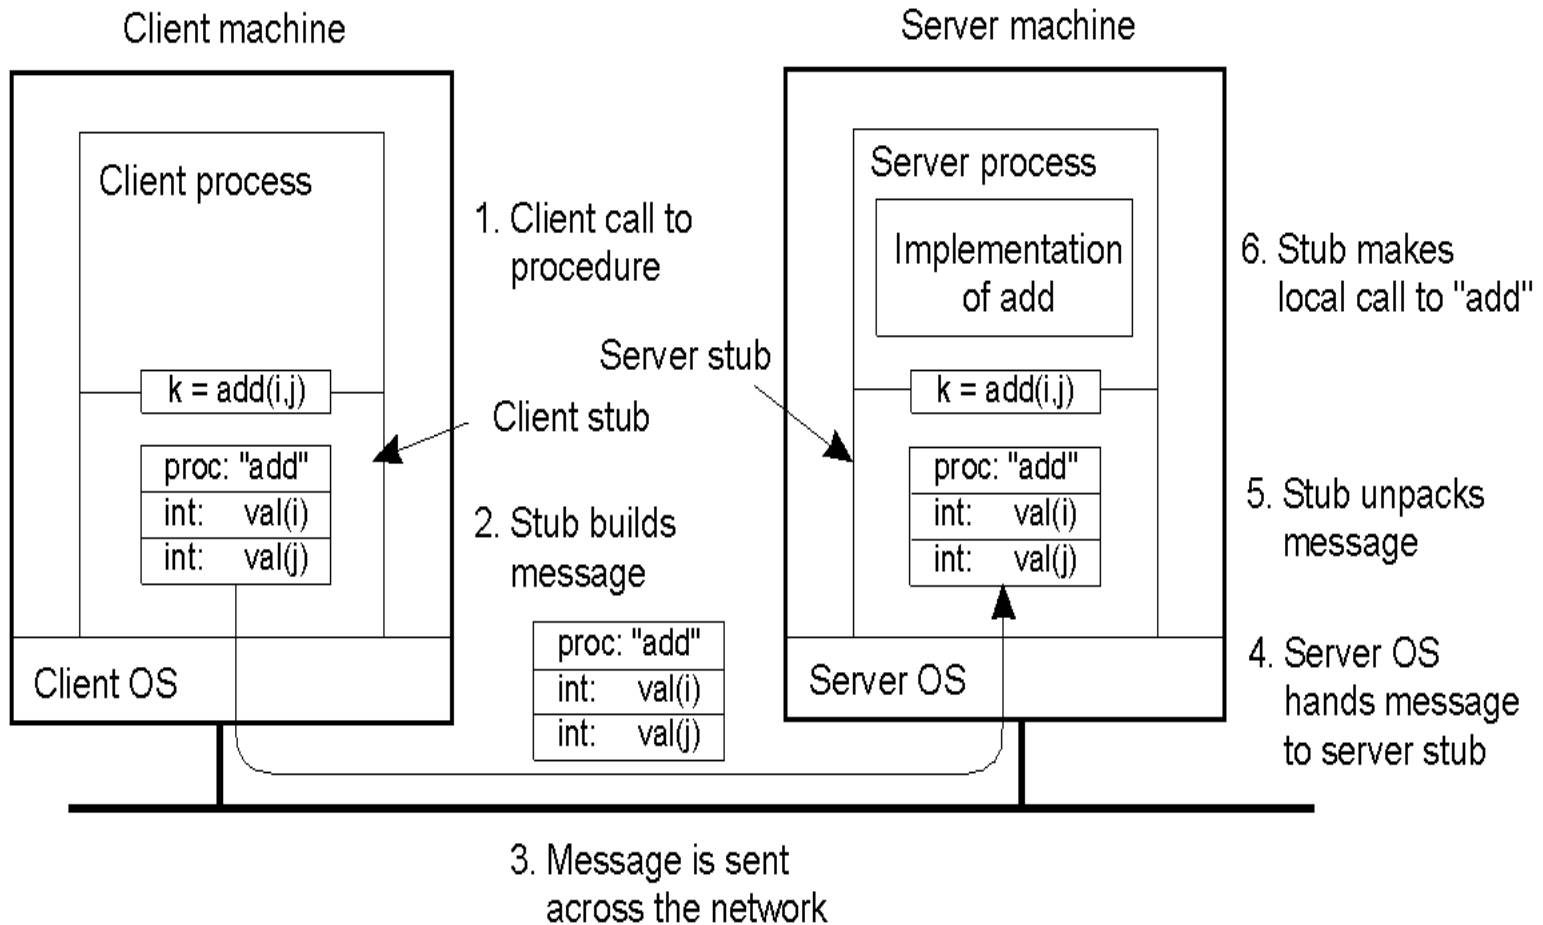
Client and Server Stubs

- Client makes procedure call (just like a local procedure call) to the client stub
- Server is written as a standard procedure
- Stubs take care of packaging arguments and sending messages
- Packaging parameters are called marshaling
- Stub compiler generates stub automatically from specs in an Interface Definition Language (IDL)
 - Simplifies programmer task

Steps of a Remote Procedure Call

1. Client procedure calls client stub in a normal way
2. Client stub builds message, calls local OS
3. Client's OS sends a message to the remote OS
4. Remote OS gives a message to the server stub
5. Server stub unpacks parameters, calls server
6. Server does work, returns result to the stub
7. Server stub packs it in a message, calls local OS
8. Server's OS sends a message to the client's OS
9. Client's OS gives a message to the client stub
10. Stub unpacks result, returns to client

Example



Steps involved in doing remote computation through RPC

RPC Challenges

- But it's not always simple
 - Calling and called procedures run on different machines, with different address spaces
 - And perhaps different environments or operating systems
 - Must convert to a local representation of data
 - Machines and networks can fail

Marshalling

- Problem: different machines have different data formats
 - Intel: little endian, SPARC: big endian
- Solution: use a standard representation
 - Example: external data representation (XDR)
- Problem: how do we pass pointers?
 - If it points to a well-defined data structure, pass a copy and the server stub passes a pointer to the local copy
- What about data structures containing pointers?
 - Prohibit
 - Chase pointers over the network
- Marshalling: transform parameters/results into a byte stream

Binding

- Problem: how does a client locate a server?
 - Use Bindings
- Server
 - Export server interface during initialization
 - Send name, version no, unique identifier, and handle (address) to binder
- Client
 - First RPC: send message to the binder to import server interface
 - Binder: check to see if a server has exported the interface
 - Return handle and unique identifier to client

Binding: Comments

- Exporting and importing incur overheads
- Binder can be a bottleneck
 - Use multiple binders
- Binder can do load balancing

RPC Implementation Failure Semantics

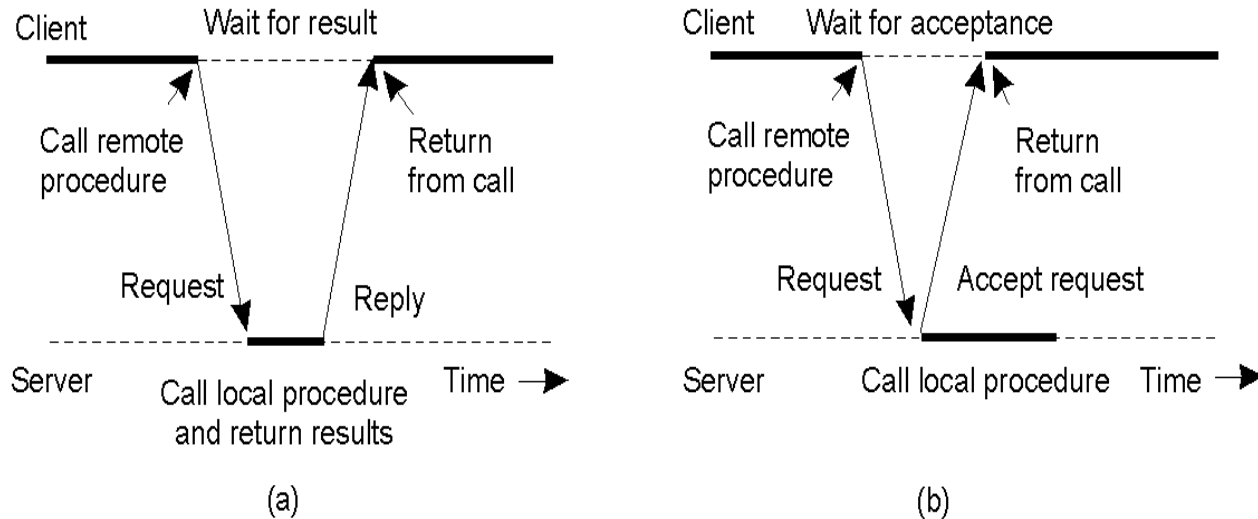
- Client unable to locate server: return error
- Lost request messages: simple timeout mechanisms
- Lost replies: timeout mechanisms
 - Make operation idempotent
 - Use sequence numbers, mark retransmissions
- Server failures: did the failure occur before or after an operation?
 - At most once
 - No guarantee

Failure Semantics

- Client failure: what happens to the server computation?
 - Referred to as an orphan
 - Extermination: log at client stub and explicitly kill orphans
 - Overhead of maintaining disk logs
 - Reincarnation: Divide time into epochs between failures and delete computations from old epochs
 - Expiration: give each RPC a fixed quantum T ; explicitly request extensions
 - Periodic checks with clients during long computations

Asynchronous RPC

- Avoids blocking of the client process.
- Allows the client to **proceed** without getting the result of the call.



- a) The interconnection between client and server in a traditional RPC
- b) The interaction using asynchronous RPC

References

- Above the Clouds: A Berkeley View of Cloud Computing, <https://www2.eecs.berkeley.edu/Pubs/TechRpts/2009/EECS-2009-28.pdf>
- Andrew Birrell and Bruce Nelson, Implementing RPCs, ACM Transactions on Computer Systems, Vol. 2, No. 1, Pages 39-59, February 1984.
- B. Bershad, T. Anderson, E. Lazowska, and H. Levy, Lightweight Remote Procedure Call, Proceedings of the 12th ACM Symposium on Operating Systems Principles, Operating Systems Review, Vol. 23, No. 5, Pages 12-113, December 1989
- Distributed Systems: Principles and Paradigms by Tanenbaum and van Steen, Chapter 4.1~4.2.